



Lab SAD2022

Queue Service on a PaaS





Motivation

- ▶ Cloud services should be designed combining microservices
- ▶ Interaction among microservices should be reactive
 - ▶ Events as messages sent among microservices
- ▶ Queues are used to store messages by microservices
 - ▶ To be consumed later on by other microservices
- ▶ A well understood pattern
 - ▶ Kafka
 - ▶ NATS
 - ▶ Artemis
 - ▶ ...
- ▶ Example application: FaaS



- ▶ Generic PUB/SUB pattern
 - ▶ Events/Messages are published to a queue by PRODUCERS
 - ▶ Messages have a topic
 - ▶ CONSUMERS express interest on one or various topics
 - ▶ Variations on how many times a message is delivered to consumers
 1. At most once (no repeat, demonstrably best effort when no failures occur)
 2. At least once
 3. Exactly once
- ▶ To increase reliability the queue process ought to be replicated
 - ▶ Also for scalability

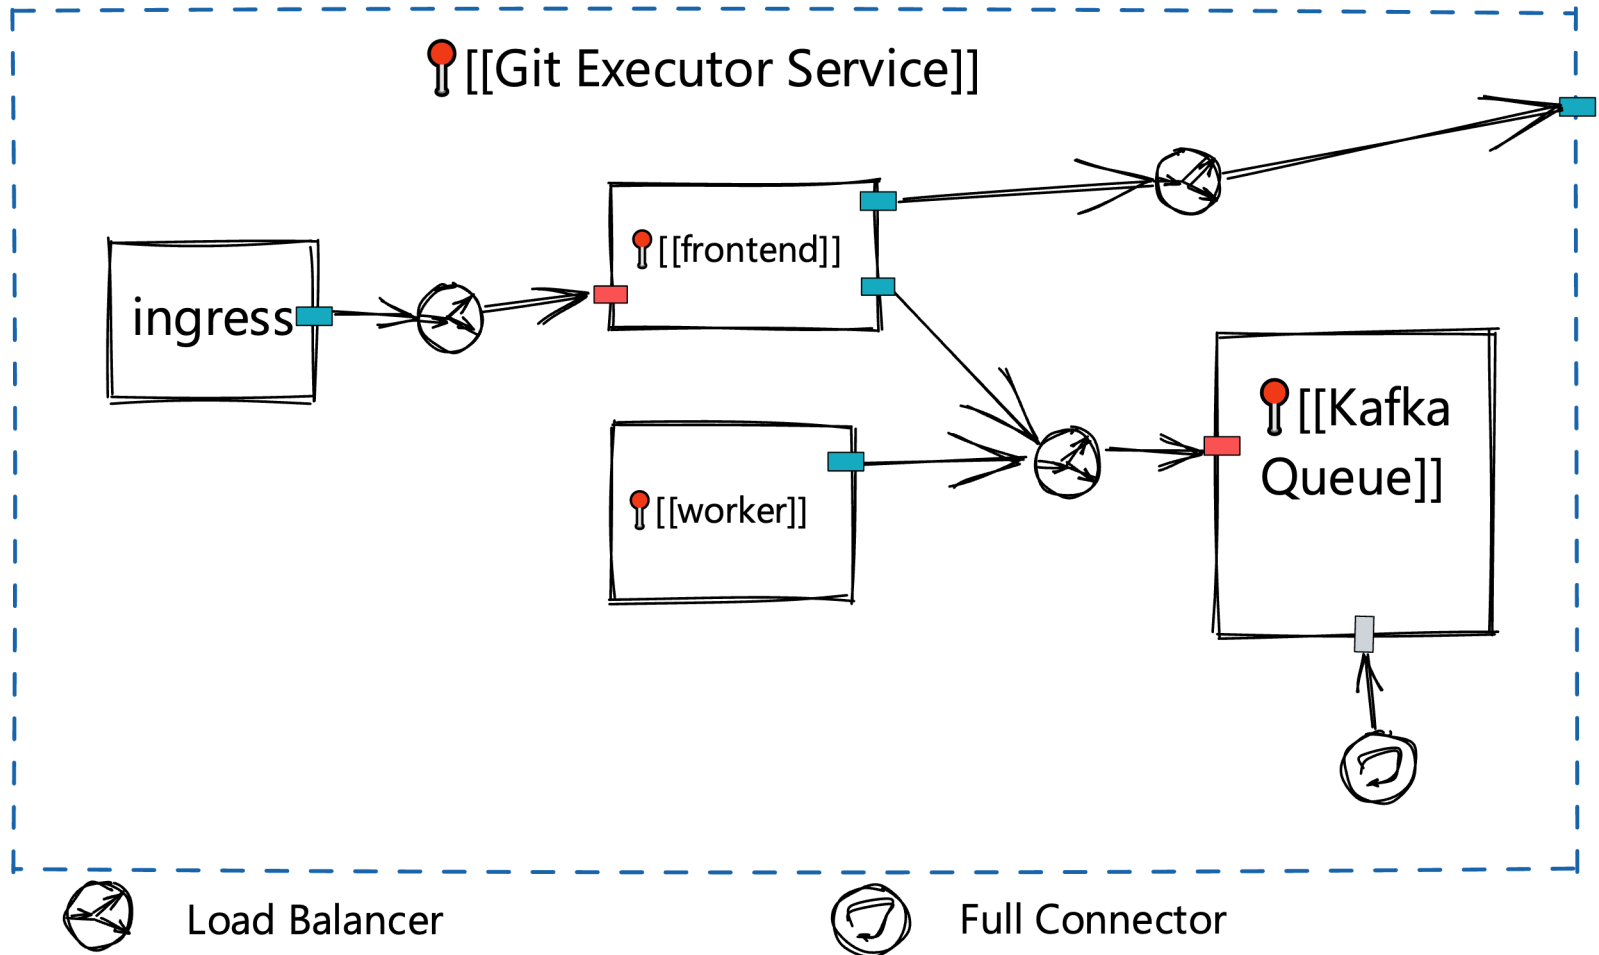


- ▶ Also called serverless
- ▶ In contraposition to SaaS: no state is maintained after a computation is requested
 - ▶ Fire and forget:
 - ▶ Request execution of a function
 - ▶ Get the results
- ▶ Implemented as a combination
 - ▶ Frontend/adaptor for the FaaS API (typically REST)
 - ▶ Queue: persist the job orders until the function executor is ready to run it
 - ▶ Job Workers, potentially specialized, capable of running the job orders and produce results



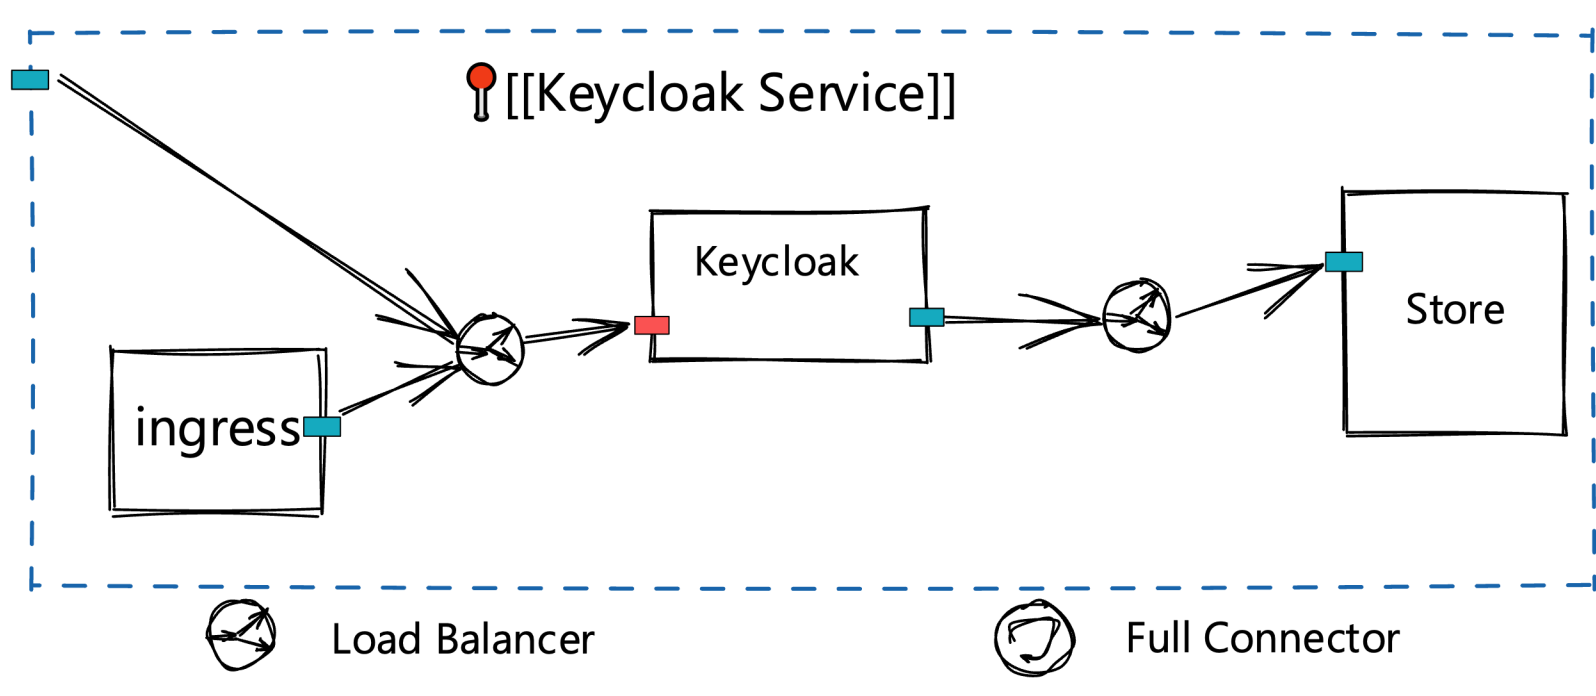
FaaS: Git Executor workers

- ▶ The complete service will have at least three microservices
 - ▶ **Frontend**
 - ▶ Receives requests from clients
 - ▶ Requests are REST, with json payload
 - ▶ Prepares them to be placed on the queue
 - ▶ **Worker**
 - ▶ Picks up work orders from a queue, carries out the job, and places the result back to another queue
 - ▶ **Queue**
 - ▶ Implements the queue functionality
 - ▶ Can coordinate with other copies of itself



FaaS: Letting Jobs in

- ▶ The FrontEnd must protect the service
- ▶ Should only let in those requests that are authenticated
 - ▶ Needs to trust an authenticator service
 - ▶ Implemented with Keycloak

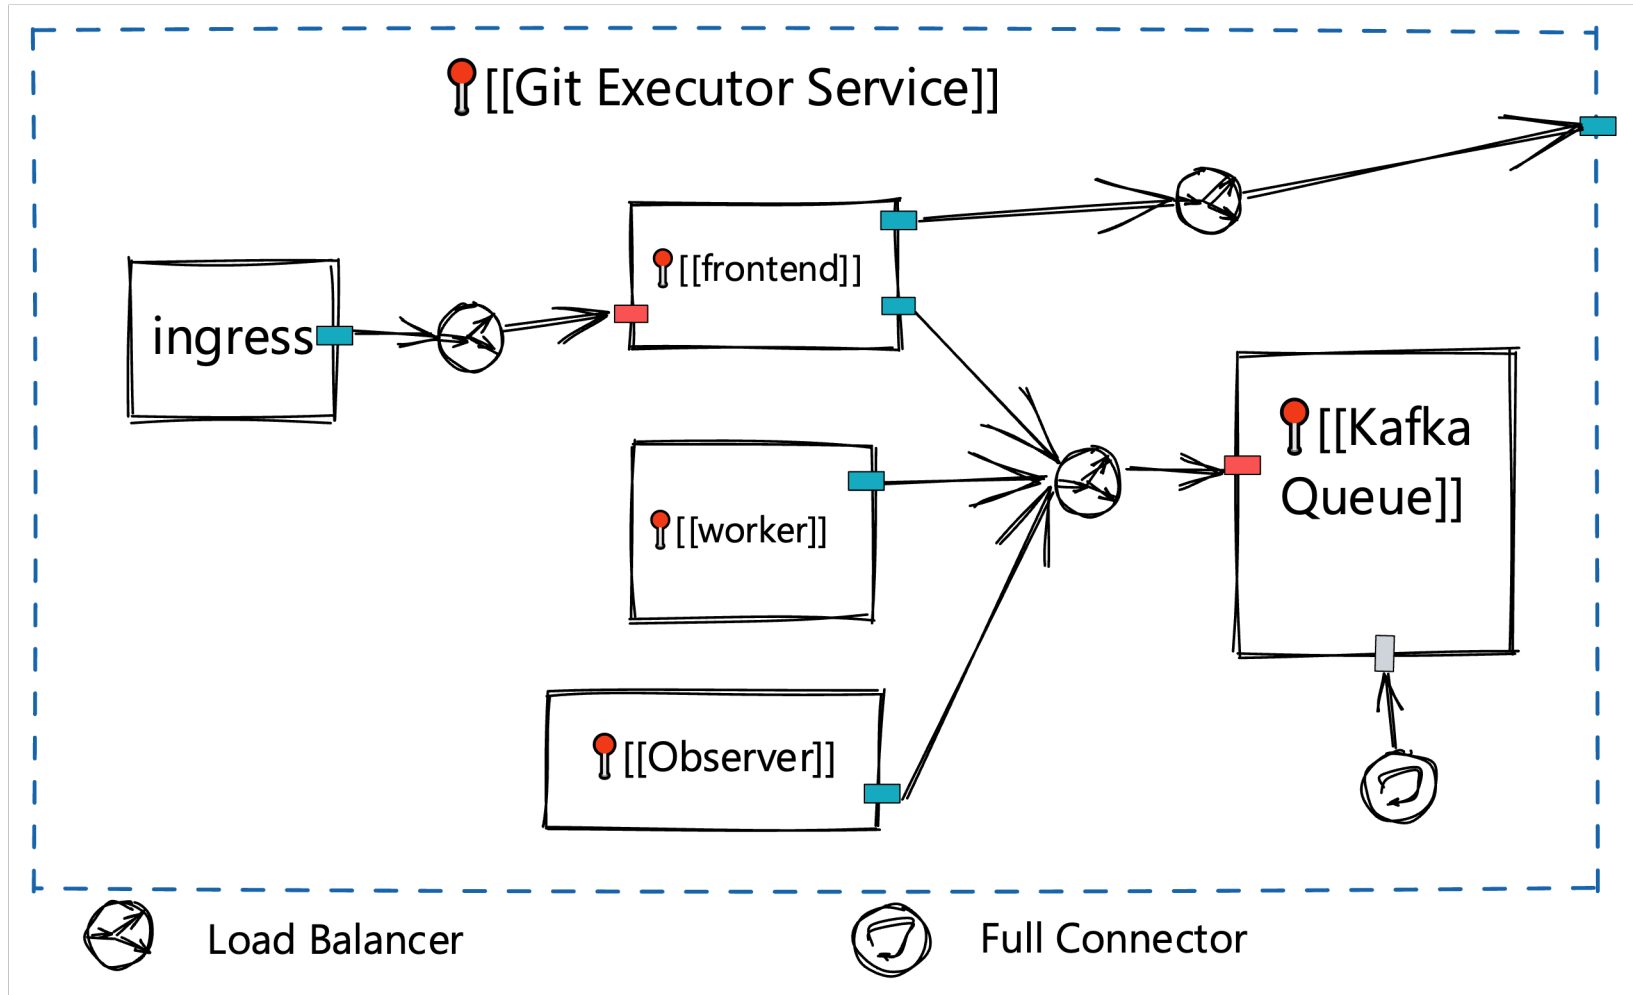




FaaS: Observer

- ▶ Microservice that inspects the queues
 - ▶ Job queue
 - ▶ Responses queue
- ▶ Propose an algorithm for elasticity
 - ▶ Based on queue length/response times
 - ▶ Configured max rate/response time
 - ▶ Determines when a new worker should be added
 - ▶ Determines when a worker must be removed
 - ▶ Sends its decisions to another queue in the queue server
 - ▶ Observed rate of arrival
 - ▶ Observed rate of completion
 - ▶ Response times

Basic Architecture with Observer





Components/Microservices

- ▶ Three components (at least)
 - ▶ Frontend, with two “channels”
 - ▶ endpoint
 - Suggestion: Implement with express or similar
 - To easily handle REST requests
 - To prepare to serve SPA client app
 - ▶ Queueclient
 - Protocol depends on Queue technology
 - ▶ Executor
 - ▶ Queueclient
 - ▶ Kafka Queue
 - ▶ discovery, coordination
 - ▶ Observer (Optional)



Implementation: Frontend

- ▶ Rest server (http)
 - ▶ Protected with Bearer Token
 - ▶ With claims about identity
- ▶ API
 - ▶ Send job/return job ID
 - ▶ Frontend injects identity of request (from bearer token)
 - ▶ Question status of job ID
 - ▶ Get result of JOB id
 - ▶ With elapsed time
 - Erases Job from queue (WARNING: observer)
 - ▶ Get list of own jobs submitted
 - ▶ Set parameters for observer (optional)
 - ▶ Subscribe to changes in jobs via websockets (optional)



Implementation: Executor

- ▶ Picks up Jobs from job queue
- ▶ Job Description:
 - ▶ GIT REPO
 - ▶ Data source links
 - ▶ Data
 - ▶ Data sink links
 - ▶ Credentials needed
 - ▶ Should be sent encrypted
 - ▶ Encryption/decription key passed in environment
 - ▶ Proceeds by cloning repo and then
 - ▶ Executing a “main” function passing data/credential parameters
 - ▶ Sending “results” to a results queue
 - ▶ Tagging result with elapsed time




Implementation: Observer

- ▶ Reads the Job Queue
- ▶ Reads the results queue
- ▶ Execs its algorithm
 - ▶ Determines info
 - ▶ Average rate of arrival
 - ▶ Average rate of service
 - ▶ Average response time
 - ▶ Add/remove/nothing worker decision
- ▶ Reports back on the observer queue



Deployment

- ▶ Each microservice's code will be encapsulated within its own docker image
- ▶ Two services linked





Deployment

- ▶ START with two docker-compose
 - ▶ One for Faas, another for keycloak
- ▶ THEN, model it within a specific PaaS:
 - ▶ <https://docs.kumori.systems>
 - ▶ Two services: FaaS/Keycloak
 - ▶ Choose any store for keycloak
 - ▶ May be in the same container
 - ▶ The two services modeled as service applications
 - ▶ The earlier images provide the architecture of the service apps.



Deliverables (yours)

- ▶ Paper, describing the architecture and functioning of the service
 - ▶ Extras you have undertaken must be included
 - ▶ Extras you think may be worth exploring
- ▶ Code
 - ▶ Folder per microservice
 - ▶ Including the Dockerfile, if needed
 - ▶ Including docker-compose files
 - ▶ Including KPaaS model manifests
 - ▶ Folder with service app manifest per service
 - ▶ Test suite used to verify the properties of your service (as described in the paper)
- ▶ In git.upv.com
 - ▶ Group: SAD2122, under SAD. You will be assigned a user per group
- ▶ DEADLINE: Jan 31



Grading

- ▶ **Base:**
 - ▶ Spec for job/response structure
 - ▶ Keycloak configuration
 - ▶ Implementations frontend/worker
 - ▶ Docker-compose files
 - ▶ Model manifests
 - ▶ Well-written paper (Spanish/english) explaining the work
- ▶ **Test code with good coverage**
- ▶ **EXTRA:**
 - ▶ Observer



Resources

- ▶ Virtual Machine in portal-ng.dsic.upv.es
 - ▶ To use docker
 - ▶ To carry out initial experiments with docker-compose
 - ▶ To use development environment of KPaaS
- ▶ Access to PaaS setup to experiment
 - ▶ Will be provided later on
- ▶ KPaaS documentation: <https://docs.kumori.systems>
- ▶ Use one of these queue systems:
 - ▶ Kafka: <https://kafka.apache.org/>
- ▶ NOTE: as much as possible, try to use the official/available docker images



Platform

- ▶ URL: admission-forge.vera.kumori.cloud
- ▶ User: your student username preceded with SAD_
 - ▶ Password...TBD



Brief introduction to Kumori PaaS

- ▶ Built on top of Kubernetes
 - ▶ Does not expose Kubernetes concepts
- ▶ Manages deployment of service applications
 - ▶ Built according to the Kumori Service Model
 - ▶ Using the Platform tools to deploy
 - ▶ Which, in turn, exercise the platform's API



Kumori's Service Model

- ▶ Set of concepts and a language to express them
- ▶ Main concepts:
 - ▶ Component
 - ▶ Service Application
- ▶ Component == the “program” of a microservice.
- ▶ Service Application == the “program” of a complete complex service
- ▶ In both cases
 - ▶ Specification of a service



Kumori's Service Model

- ▶ Specification of a service
 - ▶ Common to components and service apps
 - ▶ Communication channels
 - ▶ Server
 - ▶ Client
 - ▶ Duplex
 - ▶ Configuration
 - ▶ Parameters
 - By value
 - ▶ Resources
 - By reference



Component-specific elements

- ▶ Vertical size
 - ▶ CPU, Memory, Bandwidth
- ▶ Code
 - ▶ Docker image
 - ▶ Or images, as more than one container can be specified
 - ▶ Mappings
 - ▶ Relates configuration parameters to elements within a container
 - Files
 - Directories
 - Environment variables
 - ▶ Starting command
 - ▶ Service discovery convention
 - ▶ DNS-based
 - ▶ Resolving client channel names



Service-app specific elements

▶ Role

- ▶ Implemented by a Component
- ▶ Equivalent of a microservice's deployment
- ▶ Multiple instances of a Role can exist in the deployment of a service.
- ▶ Specifies how the configuration of the application transforms into the configuration of the underlying component
- ▶ Specifies "horizontal size" properties
 - ▶ Instances
 - ▶ Resilience
- ▶ Inherits the channels of its component



Service-app specific elements

▶ Connector

▶ Intermediates communications between roles

- ▶ Connects client channels to server channels
- ▶ Also interconnects duplex channels

▶ Two kinds:

▶ LB (load balancer)

- From client to server
- Chooses the instance of the server that receives the connection
- PURPOSE: Client-server patterns

▶ Complete/Full


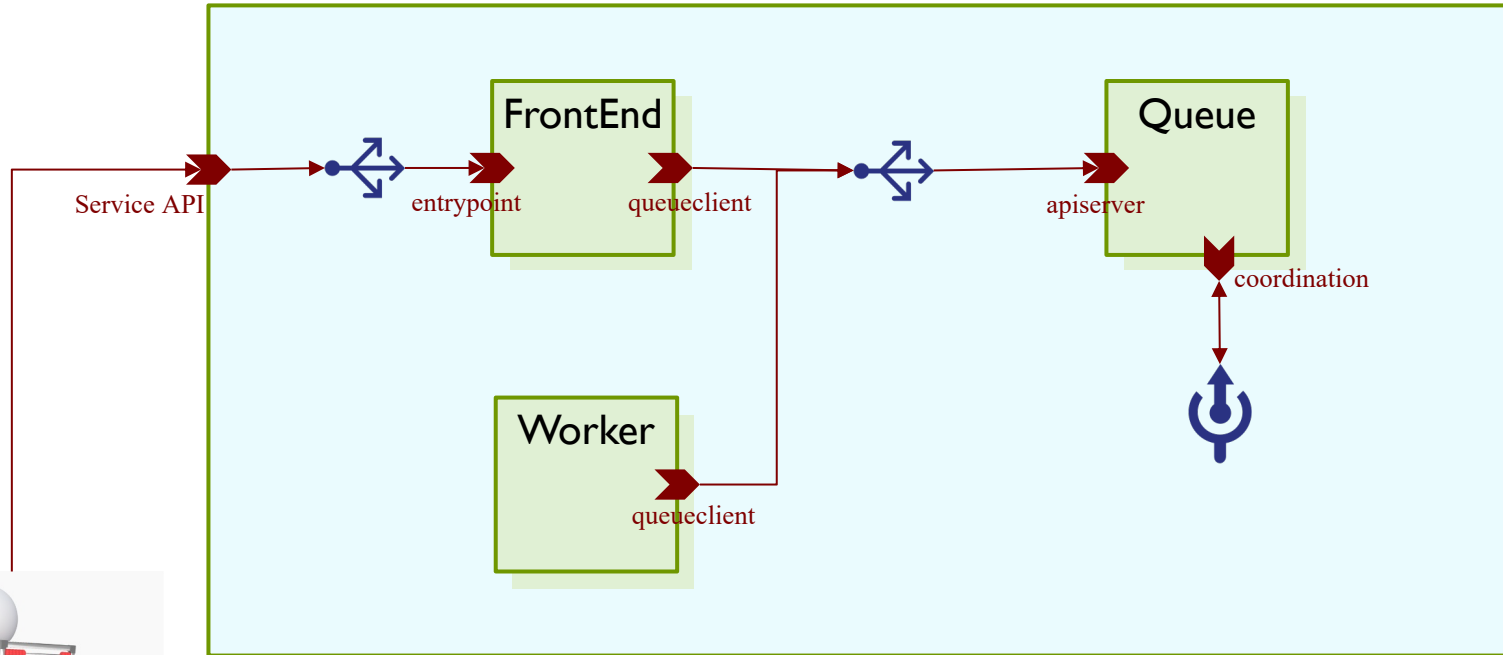
- To duplex
- From client to server
- Exposes to client (or duplex) all instances of server (duplex) channels.
- PURPOSE: Peer to Peer pattern



Service-app specific elements

- ▶ **Links**
 - ▶ From client channels in roles to connectors
 - ▶ Or
 - ▶ From connectors to server (or duplex) channels in roles
 - ▶ Only complete connectors can be linked to duplex channels
- ▶ **The topology of the service is the graph resulting from**
 - ▶ Roles
 - ▶ Connectors
 - ▶ Links
- ▶ **Graphical representation**
 - ▶ Server channels on the left
 - ▶ Client channels on the right
 - ▶ Duplex channels at the bottom

Example service application specification



 Load balancer connector

 Complete connector



Formalism

- ▶ Components and Service apps speced in CUE
 - ▶ <https://cuelang.org>
 - ▶ Gaining adoption within the devops ecosystem
- ▶ Syntactically a superset of JSON
 - ▶ All JSON is syntactically CUE
- ▶ Logical language with two main operations
 - ▶ Disjunction
 - ▶ Unification
- ▶ Useful built-in operators



Tooling

- ▶ An instance of the platform can be
 - ▶ A cluster in a private CPD
 - ▶ Physical
 - ▶ Virtual
 - ▶ Hybrid
 - ▶ A set of VMs in an IaaS provider
 - ▶ A mixed environment
- ▶ kam
 - ▶ CLI tool used in development and deployment
 - ▶ Interacting with instances of the platform
 - ▶ Needs Node.js installed.
 - ▶ Installed via npm: `npm install -g @kumori/kam`
 - ▶ Expect to update frequently
 - ▶ General format
 - ▶ `kam <subcommand> <options and arguments>`



Tooling: kam

- ▶ Authoring kpaas modules
 - ▶ `kam mod init <module name>`
 - ▶ `kam mod dependency <dependent module name>`
- ▶ Deploying services
 - ▶ `kam ctl deploy <deployment manifest path>`
- ▶ Setting up a context cluster
 - ▶ `kam ctl config -a <admission url for cluster>`
 - ▶ `kam ctl login <username>`
- ▶ Obtaining list of registered resources
 - ▶ `kam ctl get resources`
 - ▶ `kam ctl get resource <resource name>`
 - ▶ `kam ctl register resource <resource name> -d <data>`
 - ▶ ...