

Distributed Systems. Applications and Services

SAD



Goals of this session

- ▶ Get an understanding of
 - ▶ What is a distributed system
 - ▶ What are the main difficulties
- ▶ Additionally
 - ▶ Explore some examples
 - ▶ Focus on Software as a means to obtain a service.



1. Concept of Distributed System
2. Goals
3. Applications vs Services
4. Application Areas



I. What is a Distributed System

- ▶ Set of autonomous agents
 - ▶ Each agent is a sequential process, proceeding at its own pace.
 - ▶ Think of it as a State Machine/Automaton
 - Its specification is the “program”
 - Input Alphabet – Output Alphabet (e.g. look at I/O Automata)
- ▶ Agents interact
 - ▶ Options:
 - ▶ Message passing
 - ▶ Shared memory (Concurrent Systems)
 - Potentially different from arbitrary distributed systems
 - Special hardware coordination primitives (e.g. test-and-set)
 - ▶ Event-driven interaction
 - **Reactive** systems




I. What is a Distributed System

- ▶ Agents have their own state
 - ▶ Independent of each other
 - ▶ Built as a result of the “reactions” to events processed.
 - ▶ The Communication Medium (e.g. Network) Also has its own state
 - ▶ Its an agent of sorts...
- ▶ There is some collective goal to this cooperation
 - ▶ By which the behavior of the whole “system” can be assessed.
 - ▶ Usually as a collection of “events”
 - Happening through each one of the elements of the system
 - Potentially enriched with state transitions at each agent

I. What is a Distributed System

A Behavior:

$e_1 \rightarrow e_2 \rightarrow e_3 \dots \rightarrow e_m \rightarrow$





I. What is a Distributed System

- ▶ Correctness of a distributed system is declared by
 - ▶ Properties of behaviors that may appear in the system
 - ▶ E.g. Only some sets of behaviors are declared “Legal”
- ▶ Behaviors from different points of view:
 - ▶ Each process
 - ▶ A subset of the processes
 - ▶ A subset of the events
 - ▶ By kind of event
 - E.g. data access messages



Index

1. Concept of Distributed System
2. Goals
3. Applications vs Services
4. Application Areas



2. Goals

- ▶ Evolving field since its beginnings
 - ▶ Offshoot of concurrent systems
 - ▶ Heavily studied for their usefulness in the design of time sharing systems
 - ▶ You should be familiar with many aspects of concurrent systems (CSD, SO, TSR)
- ▶ Pushed by evolution of computer networks
 - ▶ How to make all those computers do something globally useful?



2. Goals

- ▶ Main original “reasons”/ “goals”
 - ▶ Speed up
 - ▶ Take a complex problem, split it in pieces, have each piece taken care of by a different computer.
 - ▶ Scale-up
 - ▶ As needs increase, add more elements to the systems and cope with the size
 - ▶ Fault tolerance
 - ▶ Basic idea:
 - If one computer breaks down, we still have other computers capable of carrying the tasks of the broken down one.
 - ▶ Resource Sharing
 - ▶ One computer may have resources (e.g., printers, disks, ...) other computers do not have (and do not need to have)
 - ▶ It should be possible to access resources from everywhere




Index

1. Concept of Distributed System
2. Goals
3. Applications vs Services
4. Application Areas

3. Concept of Service

- ▶ Specification of how an interaction should proceed
 - ▶ What “behaviors” are acceptable



- ▶ Familiar concept... from what a distributed system model says
$$e_1 \rightarrow e_2 \rightarrow \dots \rightarrow e_n \rightarrow \dots$$
$$R \rightarrow R' \rightarrow r \rightarrow \dots \rightarrow r' \rightarrow \dots$$
- ▶ Restrictions are placed on the events emitted by the clients



3. Concept of Application

- ▶ Executable specification of how a system should react to events
 - ▶ In an unambiguous language
 - ▶ Programming languages
 - ▶ An “interpreter” can execute it
 - ▶ Sets up an environment
 - To receive events and ...
 - ... process those events according to the spec ...
 - ... causing reactions
 - state changes
 - Effects for external entities (e.g., messages)
 - ▶ STATE is maintained for a running application
- ▶ A behavior can be established for a running application
 - ▶ According to the pattern of arriving events
 - ▶ *If that pattern satisfies X then the resulting behavior satisfies Y*



3.Applications and Services

- ▶ Services are obtained from the execution of an application
- ▶ Services must satisfy restrictions on their behaviors
 - ▶ That is their Service Level Agreement – SLA
 - ▶ At least two aspects
 - ▶ Functional
 - Safety properties
 - ▶ Performance
 - Liveness properties



Index

1. Concept of Distributed System
2. Goals
3. Applications vs Services
4. Application Areas



4. Application Areas: The WWW

- ▶ Servers wait for simple requests for documents
 - ▶ Requests may involve reading or writing of a document
 - ▶ But that logic is internal to the server
- ▶ The Clients are Browsers, sending/receiving documents
 - ▶ Browsers parse documents searching for metadata
 - ▶ Browsers do not need to be “Browsers” for all applications
 - ▶ Links are particular metadata pointing to other documents
 - ▶ Documents may be in different servers
- ▶ Simple and powerful paradigm
 - ▶ Initially conceived for document sharing
 - ▶ Extended to allow document requests to stand for general service requests
 - ▶ Returned “documents” encode the result of the actual request



4. Application Areas: Sensor Networks

- ▶ Driven by declining costs of hardware
- ▶ Special purpose mini-computers
 - ▶ Motes
- ▶ Embedded in common devices
 - ▶ Dishwashers, etc...
- ▶ Contain physical world sensors
 - ▶ Humidity, temperature, power consumption, ...
- ▶ Wide range of potential applications
 - ▶ Surveillance
 - ▶ Biological and chemical disaster detection
 - ▶ Power monitoring
 - ▶ ...



4. Application Areas: Cloud Computing

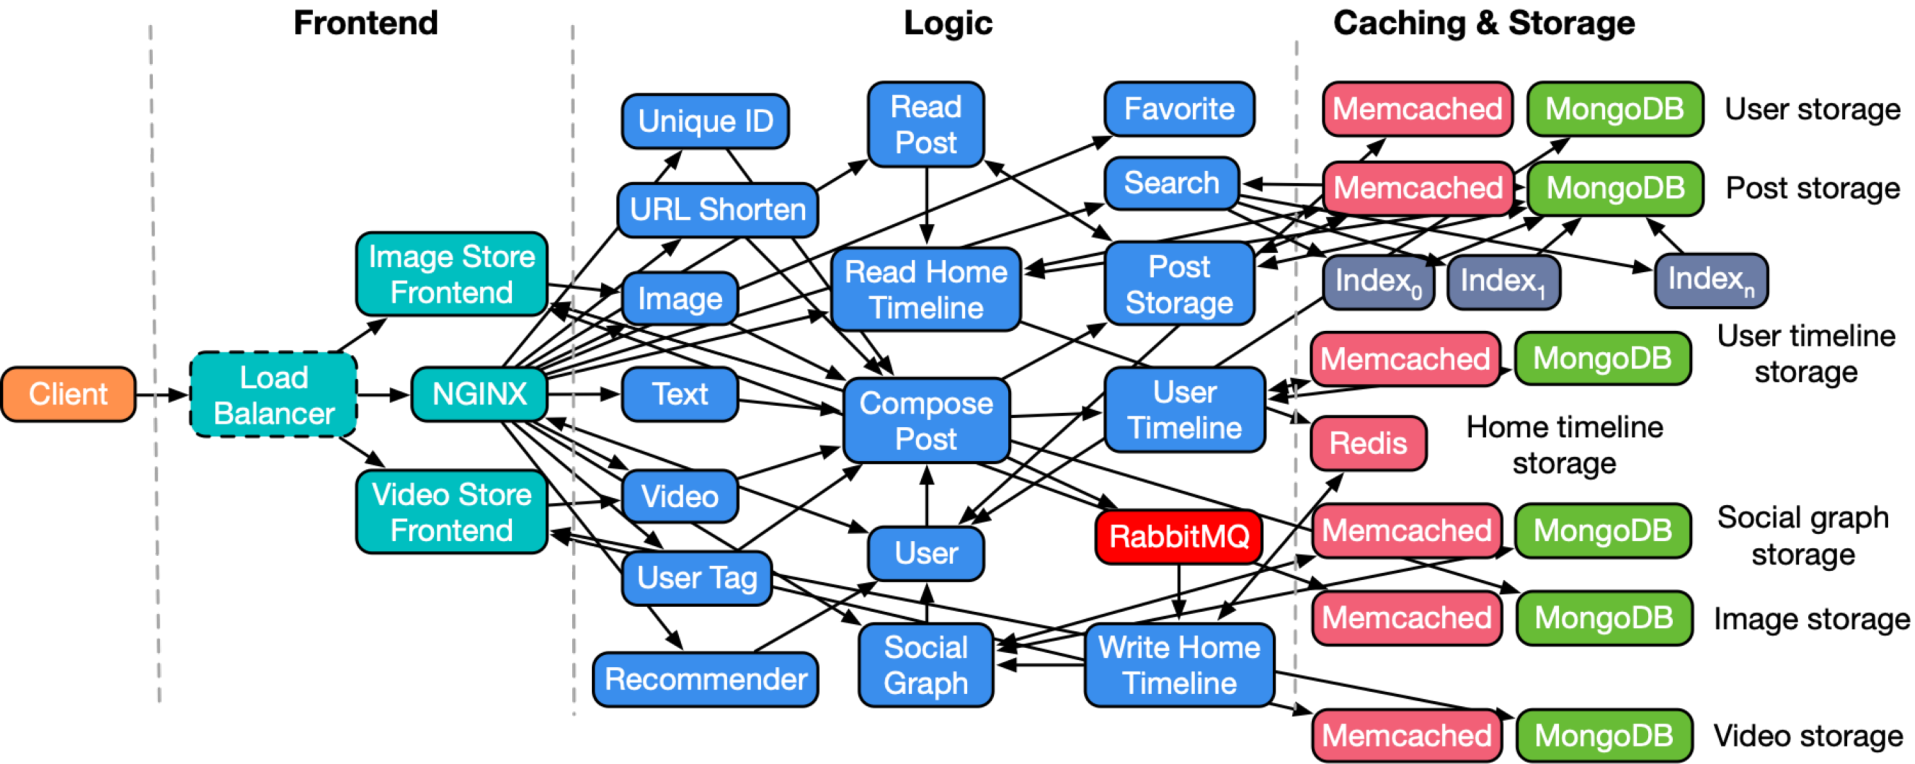
- ▶ SaaS providers
 - ▶ High volume
 - ▶ Individually
 - Google: 40,000 queries per second (1.2 Trillion per year)
 - YouTube: 1.9B active users per month,
 - viewing 5B videos per day
 - Facebook: 2.23B active users,
 - 8B video views
 - 15M photos uploaded per day
 - ▶ Combined
- ▶ At these scales, no computer can keep up.
 - ▶ By the nature of the cloud, it *has* to be massively parallel!
- ▶ IaaS providers
 - ▶ Pay-per-use



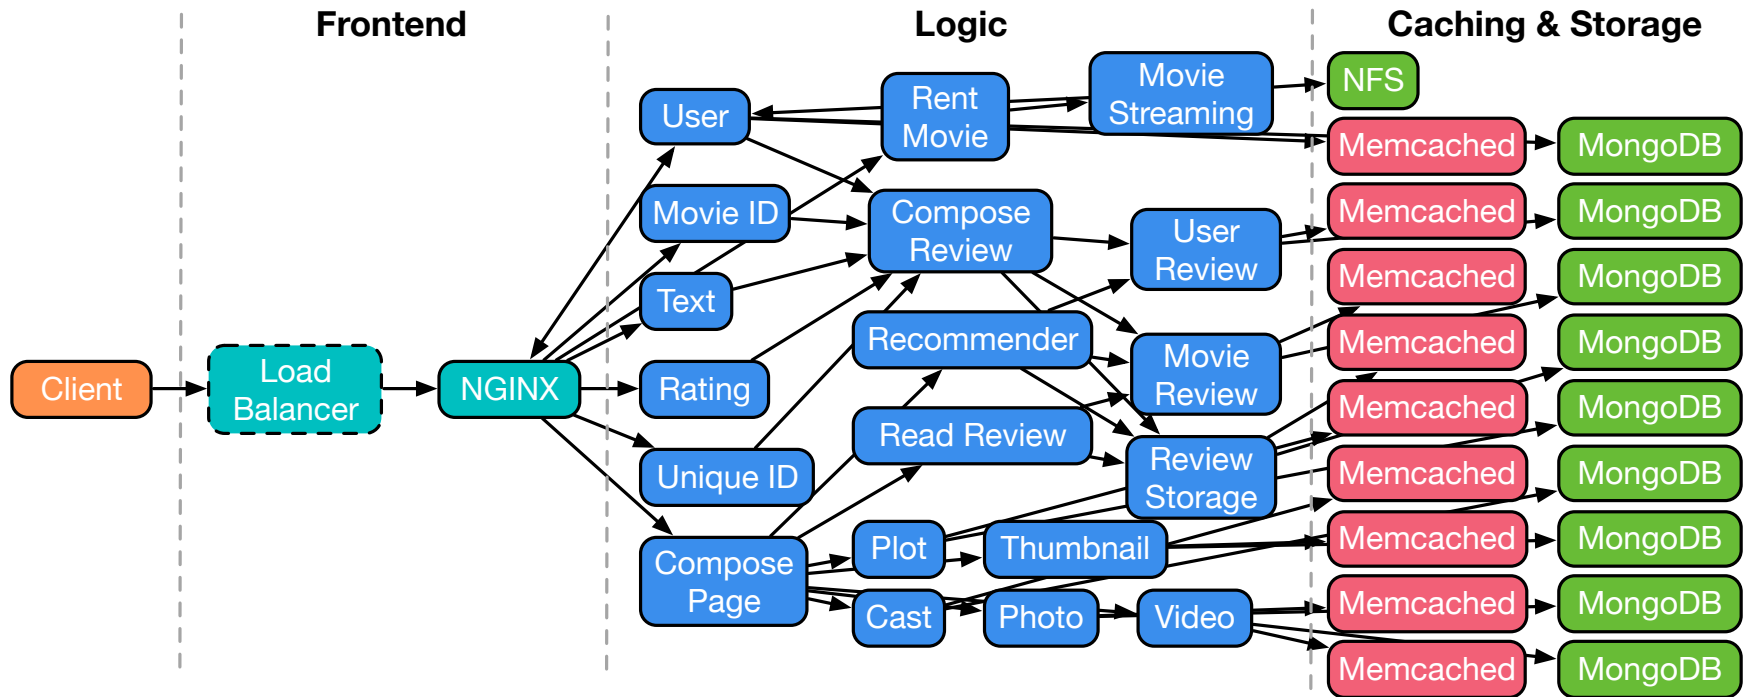
4. Application Areas: Cloud Computing

- ▶ Multi-tier
- ▶ Tier-one
 - ▶ First-line of interaction
 - ▶ Generates responses
 - ▶ Use small amounts of memory
 - ▶ Low computer power
 - ▶ Low needs for storage/network I/O
- ▶ Tier-two
 - ▶ Support services
 - ▶ Micro-services
 - Specialize on various tasks
 - E.g. Storage
 - ▶ Larger computers

4. Application Areas: Cloud/Microservices



4.Application Areas: Cloud/Microservices





4. Application Areas: Cloud/Microservices

- ▶ Every one of those little nodes is itself a small elastic pool of processes
- ▶ Microservice:
 - ▶ Out of an application
 - ▶ The data center can run one instance
 - ▶ ... or many instances, “elastically”,
 - ▶ ... to deal with dynamically varying demand.
 - ▶ Horizontal scaling
- ▶ Any instance can handle any request equally well,
 - ▶ So, no need for very careful “routing” of specific requests to specific instances.
 - ▶ This lets the data center adapt to changing loads easily!



4. Application Areas: Cloud/Microservices


- ▶ Advantages of Micro-services
 - ▶ Modular → easier to understand
 - ▶ Speed of development & deployment
 - ▶ On-demand provisioning, elasticity
 - ▶ Language/framework heterogeneity



4. Application Areas: The “Internet of Things”

- ▶ Motivation: leverage ubiquitous connectivity of all devices
 - ▶ Generalization of sensor networks
 - ▶ All devices can, and will interact among them
 - ▶ Devices can also alter their physical environment
 - ▶ New scenarios open up
 - ▶ Smart cities
 - ▶ Building/Factory automation
 - ▶ Healthcare
 - ▶ ...

The “Internet of Things”





4. Application Areas: Cooperative Computing

- ▶ Most computational power is underused
 - ▶ The desktops spend many hours doing nothing
- ▶ Many engineering and scientific problems can be split into pieces (tasks)
 - ▶ Each task can be resolved in a small amount of time
 - ▶ Results from each piece can be composed to complete the resolution of the whole problem
- ▶ Servers can be set up with an instance of such a problem
 - ▶ The server creates a pool of smaller tasks
- ▶ Computers across the internet can subscribe to receive tasks to solve
 - ▶ They install a special client software: the task runtime environment
 - ▶ The client registers with the server
- ▶ The server spreads tasks among the registered clients, and collects their results



4. Application Areas: High Availability Clusters

- ▶ So far we have seen application areas addressing resource sharing and cooperation.
- ▶ Fact:
 - ▶ Devices fail. Computers are devices. They fail at some point with a 100% probability.
- ▶ Fact:
 - ▶ Not all devices fail at the same time, always.
 - ▶ Q: when can it happen?
- ▶ Some environments need a high degree of availability
 - ▶ Banking
 - ▶ Finances
 - ▶ Mission-critical systems
 - ▶ ...
- ▶ Leverage having more than one device to stand failures



4. Application Areas: High Availability Clusters

- ▶ **HA Cluster:**
 - ▶ Set of computers, with server programs on which clients depend constantly
 - ▶ Typically holding sensitive data
 - ▶ Designed with specific protocols to stand failures of one or more of them
 - ▶ Two main concerns:
 - ▶ Preserve data integrity/consistency
 - ▶ Preserve server operation availability



4. Application Areas: Distributed Ledgers

- ▶ Avoidance of centralized authorities to keep track of transactions
- ▶ Avoidance of centralized Notarial authorities
- ▶ Technology
 - ▶ Blockchain algorithms