

# Consistent Hashing

SAD



# P2P: Overlay Networks

---

- ▶ P2P applications need to
  - ▶ Track identities & IP addresses of peers
    - ▶ May be many and may have significant churn
  - ▶ Route messages among peers
    - ▶ If you don't keep track of all peers, this is “multi-hop”
- ▶ Overlay network
  - ▶ Peers doing both naming and routing
  - ▶ IP network becomes the low level transport



# Structured Overlays

---

- ▶ Consider problem of data partition:
  - ▶ Have  $D$  resources
  - ▶ Need to store them in  $K$  servers
  - ▶ Need to distribute the load among the servers
  - ▶ Given document  $U$ , choose one of  $k$  servers to use to access it



# Consistent Hashing

---

- ▶ Initial motivation (1997)
  - ▶ Web Page Caching
    - ▶ Increase scalability of web sites
    - ▶ Avoid swamping the “home” server for the page
    - ▶ Caching is good ... but... Cache servers must be found
      - And we must avoid swamping them too
  - ▶ E.g. If we want to go to [upv.es](http://upv.es)...
    - ▶ Where do we actually go?
    - ▶ How about the browser itself?
    - ▶ If stale...
      - Why not a shared cache?
        - Difficult



# Consistent Hashing

---


- ▶ Benefits of shared Caching
  - ▶ Larger cache size
  - ▶ Higher chance of cache hit
  - ▶ Lower load on Servers
  - ▶ Faster Lookup
- ▶ BTW: Heard of Akamai? Anyone?
  - ▶ (12BUSD)
- ▶ Why is it difficult?
  - ▶ Cannot use just one machine
  - ▶ Spread over many different machines



# Consistent Hashing

---

- ▶ OK, assume you have 100 cache servers
- ▶ Approaches
  - ▶ Poll each machine in the set and ask
    - ▶ ... Expected load per server?
    - ▶ ... Is this reasonable?
  - ▶ Can we find a cache server number where to go given a URL?
    - ▶ Actually we can set up such a scheme
    - ▶ How about using a hash function?
    - ▶  $\text{Hash}(\text{url}) \rightarrow \text{Server Number where to look for the url resource}$





# Consistent Hashing

- ▶ Suppose we use modulo hashing
- ▶ Number servers  $1..k$ 
  - ▶  $X = \text{toNumber}(\text{url})$  (e.g. integer as a concat of bytes)
  - ▶ Place document with id  $X$  on server  $i = (X \bmod k)$ 
    - ▶ *Thus  $\text{hash}(\text{url}) = \text{mod}k(\text{toNumber}(\text{url}))$*
- ▶ Problem?
  - ▶ Data may not be uniformly distributed
- ▶ MOD does not have distribution guarantees
  - ▶ Distribution governed by ID space distribution
    - ▶ Which may be very non-uniform
- ▶ THUS, potential overload problems on some of the servers
  - ▶ Poor balancing





# Consistent Hashing

- ▶ A good hash function ideally must
  - ▶ Be easy to evaluate
  - ▶ Behave like a totally random function
    - ▶ Spread out universe  $U$  uniformly
- ▶ Difficult to design Good hashes
  - ▶ Better use already existing ones
  - ▶ Crypto hashes have good properties
    - ▶ MD5, SHA-1, ...
- ▶ So, assuming we have a good hash function...
  - ▶ Place url on server  $i = \text{hash}(\text{url}) \bmod k$ 
    - (Why do we need the mod?)



# Consistent Hashing

---

- ▶ Are we done?
- ▶ Assume we add a new server
  - ▶ Now we have 101 servers
- ▶ Is  $\text{hash}(u) \bmod 100 == \text{hash}(u) \bmod 101$ ?



# Consistent Hashing

- ▶ Are we done?
- ▶ Assume we add a new server
  - ▶ Now we have 101 servers
- ▶ Is  $\text{hash}(u) \bmod 100 == \text{hash}(u) \bmod 101$ ?
  - ▶ UNLIKELY
- ▶ **THUS**
  - ▶ We will need to relocate
  - ▶ **CHANGING K FORCES** resource relocation!
    - ▶ Rehash is an expensive operation
    - ▶ Disaster for applications where K changes often
      - As is the case for web page caching
        - Original motivation
      - And is the case for P2P networks



# Consistent Hashing

---

- ▶ Why  $k$  changes?
  - ▶ Failures occur
    - Thus the set of servers temporarily decreases
  - ▶ Load changes
    - It also increases to cope with the load
  - ▶ Thus → The set of servers changes over time
    - Guaranteed!
- ▶ How many do we need to relocate in the general case?

- ▶ Why  $k$  changes?
  - ▶ Failures occur
    - Thus the set of servers temporarily decreases
  - ▶ Load changes
    - It also increases to cope with the load
  - ▶ Thus → The set of servers changes over time
    - Guaranteed!
- ▶ How many do we need to relocate in the general case?
  - Even with just 1 change,  $n/(n+1) * U$  resources will need to move
    - Seems impractical for large sets



# Consistent Hashing

---

- ▶ On top of that
  - ▶ Reconfiguration takes time
  - ▶ Thus → Clients may have different views of the available set of servers
- ▶ Inconsistent view of who holds what
  - ▶ Need to maintain old state while moving
  - ▶ Need to detect when old state can be deleted
- ▶ Can we do better?
  - ▶ Hashing a good idea but...
  - ▶ We need something better
    - ▶ Enter Consistent Hashing



# Consistent Hashing

---

- ▶ What do we want?
  - ▶ To have our cake and eat it !!!
- ▶ To Have a Hash-like functionality
  - ▶ Including its attractive programming interface
- ▶ To keep most resources assigned where they were before changes
  - ▶ Thus
    - ▶ Avoid downtime on server failures
    - ▶ Avoid overloading new servers on addition to the network
    - ▶ Avoid rehashing more keys than necessary on membership changes
- ▶ Can we do it?



# Consistent Hashing

---

- ▶ **Key Idea:**
  - ▶ In addition to hashing the names of the resources...





# Consistent Hashing

---

- ▶ Key Idea:
  - ▶ In addition to hashing the names of the resources...
- ▶ **WE ALSO HASH THE Ids of the Servers!**
  - ▶ To the same range as resources
  - ▶ They become also resources of sorts




# Consistent Hashing

---

- ▶ Key Idea:
  - ▶ In addition to hashing the names of the resources...
- ▶ **WE ALSO HASH THE Ids of the Servers!**
  - ▶ To the same range as resources
  - ▶ They become also resources of sorts


So WHAT?

What is the big deal?



- ▶ Consistent hashing partitions resource key-space among “nodes”
- ▶ Contact appropriate server to lookup/store key
  - ▶ Blue nodes determines red node is responsible for key<sub>1</sub>
  - ▶ Blue node sends lookup or insert to red node


# Consistent Hashing: Common Virtual Addresses




## ► Partitioning key-space among nodes

- Nodes choose random identifiers: e.g., **hash(IP)**
- Keys randomly distributed in ID-space: e.g., **hash(URL)**
- Keys assigned to node “nearest” in ID-space
- Spreads ownership of keys “evenly” across nodes

- ▶ Each element of the array is a hash table bucket
  - ▶ Servers fall on particular buckets
  - ▶ Each resource,  $x_i$  is assigned to the server nearest to its right




- ▶ Imagine bucket array
  - ▶ Each element of the array is a hash table bucket
  - ▶ Might be very large: e.g.  $2^{32}$  for a 32 bit key length
    - ▶ Virtual, not implemented like that
  - ▶ Servers fall on particular buckets
  - ▶ Each resource,  $x_i$  is assigned to the server nearest to its right




- ▶ But what do we do with the last key?  
 $2^{32} - 1$
- ▶ What is to its right?

# Consistent Hashing

- ▶ We can glue 0 and  $2^{32} - 1$  together forming a ring
  - ▶ Servers and resources map to bucket in the ring
  - ▶ Resources are assigned to the server closest clockwise
  - ▶ Solves the problem of the last object to the right



- ▶ **N servers**
  - ▶ Partition the ring into N segments
  - ▶ Each server is responsible for one of those segments








# Consistent Hashing

---

- ▶ Nice properties
  - ▶ Assuming reasonable hash function
- ▶ By Symmetry
  - ▶ The expected load on each of the  $n$  servers is  $1/n$  fraction of the resources
    - ▶ However there is some variance we need to reduce
  - ▶ If we add a new server we only need to move those resources that need to be stored there
  - ▶ ...

# Consistent Hashing

- ▶ If we add a new server,  $s_3$ 
  - ▶  $x_2$  moves from  $s_0$  to  $s_3$
  - ▶ The rest remain at the same servers





# Consistent Hashing

---

- ▶ Nice properties
  - ▶ Assuming reasonable hash function
- ▶ By Symmetry
  - ▶ The expected load on each of the  $n$  servers is  $1/n$  fraction of the resources
    - ▶ However there is some variance we need to reduce
  - ▶ If we add a new server we only need to move those resources that need to be stored there
- ▶ Combined
  - ▶ We only need to move  $1/n$  th of the resources
  - ▶ Compare to the naïve solution
    - ▶ Only  $1/n$  th of the servers DO NOT HAVE to move

All on average...



# Consistent Hashing

---

- ▶ So, how do we implement all this efficiently?
  - ▶ **Lookup** and **Insert** operations
  - ▶ Need efficient clockwise scan operation
    - ▶ For any  $x$
    - ▶ Search for server  $s$  minimizing  $h(s)$ , such that  $h(s) \geq h(x)$
- ▶ Need data structure
  - ▶ Keys:  $h(s)$
  - ▶ Values:  $s$
  - ▶ Fast **Successor** operation
- ▶ What can we use?



# Consistent Hashing

---

- ▶ Hash table?



# Consistent Hashing

---

- ▶ Hash table?
  - ▶ NO... Does not keep order at all...



# Consistent Hashing

---

- ▶ Hash table?
  - ▶ NO... Does not keep order at all...
- ▶ Heap?



# Consistent Hashing

---

- ▶ Hash table?
  - ▶ NO... Does not keep order at all...
- ▶ Heap?
  - ▶ NO... Only keeps a partial order
    - ▶ Just to identify minimum
- ▶ Then what?





# Consistent Hashing

---

- ▶ Hash table?
  - ▶ NO... Does not keep order at all...
- ▶ Heap?
  - ▶ NO... Only keeps a partial order
    - ▶ Just to identify minimum
- ▶ Then what?
- ▶ How about a binary search tree?
  - ▶ Implements total order
  - ▶ Provides a Successor operation



# Consistent Hashing

---

- ▶ Hash table?
  - ▶ NO... Does not keep order at all...
- ▶ Heap?
  - ▶ NO... Only keeps a partial order
    - ▶ Just to identify minimum
- ▶ Then what?
- ▶ How about a binary search tree?
  - ▶ Implements total order
  - ▶ Provides a Successor operation
    - ▶  $O(\text{height of the tree})$



# Consistent Hashing

---

- ▶ Hash table?
  - ▶ NO... Does not keep order at all...
- ▶ Heap?
  - ▶ NO... Only keeps a partial order
    - ▶ Just to identify minimum
- ▶ Then what?
- ▶ How about a binary search tree?
  - ▶ Implements total order
  - ▶ Provides a Successor operation
    - ▶  $O(\text{height of the tree})$
    - ▶ Height =  $\log(n)$  when tree is balanced



# Consistent Hashing

---

- ▶ Can we reduce the variance?
  - ▶ We expect  $1/n$  th of the resources on each server on average
  - ▶ But it is very unlikely each server has that portion
    - ▶ Imagine getting a perfect partition of the ring: unlikely
- ▶ We can decrease variance by




# Consistent Hashing

---

- ▶ Can we reduce the variance?
  - ▶ We expect  $1/n$  th of the resources on each server on average
  - ▶ But it is very unlikely each server has that portion
    - ▶ Imagine getting a perfect partition of the ring: unlikely
- ▶ We can decrease variance by...
  - ▶ Making  $v$  virtual copies of each actual server
    - ▶ We can hash each one of them like this:  $\text{hash}(i,j)$ , where
      - $i$  is the identity of the server
      - $j$  in  $1..v$  is the identity of the  $j$ -th virtual server associated to server  $i$ .

# Consistent Hashing

- ▶ Suppose  $v = 4$ , for servers  $\{1,2,3\}$ 
  - ▶ We would have  $4 \times 3 = 12$  buckets in the ring, 4 per server
  - ▶ When we search for  $x$ , we do as before  $\text{hash}(x)$ 
    - ▶ Find the first bucket clockwise
    - ▶ Chose the server it belongs to





# Consistent Hashing

---

- ▶ Again, by symmetry
  - ▶ Each server is expected to get  $1/n$  th of resources on average
- ▶ But, the replication increases the number of stored keys by a factor of  $v$
- ▶ But load variance across servers is reduced significantly

- ▶ Again, by symmetry
  - ▶ Each server is expected to get  $1/n$  th of resources on average
- ▶ But, the replication increases the number of stored keys by a factor of  $v$
- ▶ But load variance across servers is reduced significantly
  - ▶ Some copies of a server will get more than  $1/vn$
  - ▶ But other will get less than  $1/vn$
  - ▶ Thus
    - ▶ We get the variance averaged out among the various copies of a server!!
- ▶ What is a good value for  $v$ ?
  - ▶ Around  $\log_2 n$  is large enough for variance
    - ▶ And small enough to avoid blowing up the size of the tree





# Consistent Hashing

---

- ▶ Additional Properties of virtualizing Servers
  - ▶ Can be brought on-line gradually
    - ▶ Start with one virtual server
      - Gradually increase the number of virtual servers it holds
  - ▶ Can have a number of virtual servers adapted to its capacity
    - ▶ The "larger" the server the more virtual servers it should have.
      - So it can take additional load
  - ▶ On failure/quitting
    - ▶ The load of the gone server is ...