

# Distributed Hash Tables

SAD



# Consistent Hashing

---

- ▶ OK.
  - ▶ This seems fine for managed clusters
  - ▶ How about P2P?
- ▶ It turns out, it can be arranged to work on non-managed sets of servers
  - ▶ Need to find a way to implement the successor operation
    - ▶ When nobody keeps track of the full set of servers
    - ▶ In a fully distributed system
    - ▶ Distributed Hash Tables (DHT)
- ▶ Exercise: Study the Dynamo System from Amazon.
  - ▶ Describe the concrete way CH is implemented there
  - ▶ Describe performance characteristics for that implementation
  - ▶ Prepare to implement a managed consistent hashing cluster



# From Consistent Hashing to DHT: Chord

---

- ▶ What is Chord?
  - ▶ A P2P lookup service
    - ▶ Find where to go for a resource
  - ▶ Problem: Locate a data item in a collection of distributed nodes
    - ▶ Frequent node arrivals and departures
- ▶ Reminder:
  - ▶ Efficient location of resource is core to many P2P systems
- ▶ Just one main operation:
  - ▶ Map a key to a node



# DHT: Chord

---

- ▶ Chord goodness?
  - ▶ Simplicity
  - ▶ Provable correctness
  - ▶ Small storage overhead for search purposes
    - ▶ Each Chord node needs routing information about a small set of other nodes
  - ▶ Maintains routing information dynamically
    - ▶ As nodes join and leave the system
- ▶ Lookups by means of messages to other nodes
  - ▶ Iteratively/recursively



# DHT: Chord

---

- ▶ NOTE: Mapping to Nodes, not Values
  - ▶ Traditional storage maps keys to values
- ▶ Easy to do in Chord
  - ▶ Each (key, value) pair can be stored at the key's node
- ▶ Depends on the application
  - ▶ Chord is multi-purpose
    - ▶ Only location of node responsible for a key
  - ▶ Application can store a key-value store on each node
    - ▶ Finally resolving the value for a key, after finding the node where the key-value is stored



# DHT: Chord

---

- ▶ We saw other approaches:
- ▶ Napster
  - ▶ Centralized Directory
    - ▶ SPoF/Dictatorship
- ▶ Gnutella et al
  - ▶ Flooding on an unstructured overlay network
    - ▶ Difficult scaling proposition
- ▶ Consistent-hashing based DHT (like Chord)
  - ▶ Scales well
  - ▶ Fully distributed: independent of any single point of control or failure



# DHT: Chord

---

## ▶ Let us compare to DNS

### DNS

- ▶ provides a host name to IP address mapping
- ▶ relies on a set of special root servers
- ▶ names reflect administrative boundaries
- ▶ is specialized to finding named hosts or services

### Chord

- ▶ can provide same service: Name = key, value = IP
- ▶ requires no special servers
- ▶ imposes no naming structure
- ▶ can also be used to find data objects that are not tied to certain machines



# DHT: Chord

---

## ▶ Addressed difficulties

### ▶ Load Balancing

- ▶ Spreading keys efficiently among nodes
  - Consistent hashing

### ▶ Decentralization

- ▶ Fully distributed algorithm
  - Non-managed

### ▶ Scalability

- ▶ Logarithmic growth of lookups (with number of nodes in network)
  - Feasible for large systems





## ▶ Addressed difficulties (cont)


### ▶ Availability

- ▶ Automatic adjustment of internal tables
- ▶ Goal: Ensure node responsible for a key is always found

### ▶ Flexible naming

- ▶ No structure imposed on keys
  - Key space is flat
- ▶ Mapping application “names” to Chord keys is up to the application
  - Chord is just enabling middleware

# DHT: Chord: Example



- ▶ Highest layer provides a filesystem-like interface
  - ▶ User-friendly names, authentication, etc
- ▶ Middle layer converts to block operations
- ▶ Block storage uses Chord
  - ▶ Identify node responsible for storing block
    - ▶ Then Block Store can talk to storage server on the holding node



# DHT: Chord

---

- ▶ Base protocol main tasks:
  - ▶ Specify how to find the locations of keys
    - ▶ E.g., how lookup, the main operation, is performed
  - ▶ Specify how to join new nodes to the system
    - ▶ How the structures are updated to maintain correct lookup
  - ▶ Specify what to do on planned node departures
    - ▶ Same as before: how are structures modified
  - ▶ Specify what to do on failures
    - ▶ Unplanned departures



# DHT: Chord

---

- ▶ Consistent hashing is used as a base, with its properties



# DHT: Chord

- ▶ Consistent hashing is used as a base, with its properties
- ▶ SUMMARY of benefits
  - ▶ Hash function assigns each node *and* key an m-bit *identifier* using a base hash function such as SHA-1
    - ▶  $ID(\text{node}) = \text{hash}(\text{IP}, \text{Port})$
    - ▶  $ID(\text{key}) = \text{hash}(\text{key})$
- ▶ Properties of consistent hashing:
  - ▶ Function balances load: all nodes receive roughly the same number of keys – good?
  - ▶ When an Nth node joins (or leaves) the network, only an  $O(1/N)$  fraction of the keys are moved to a different location

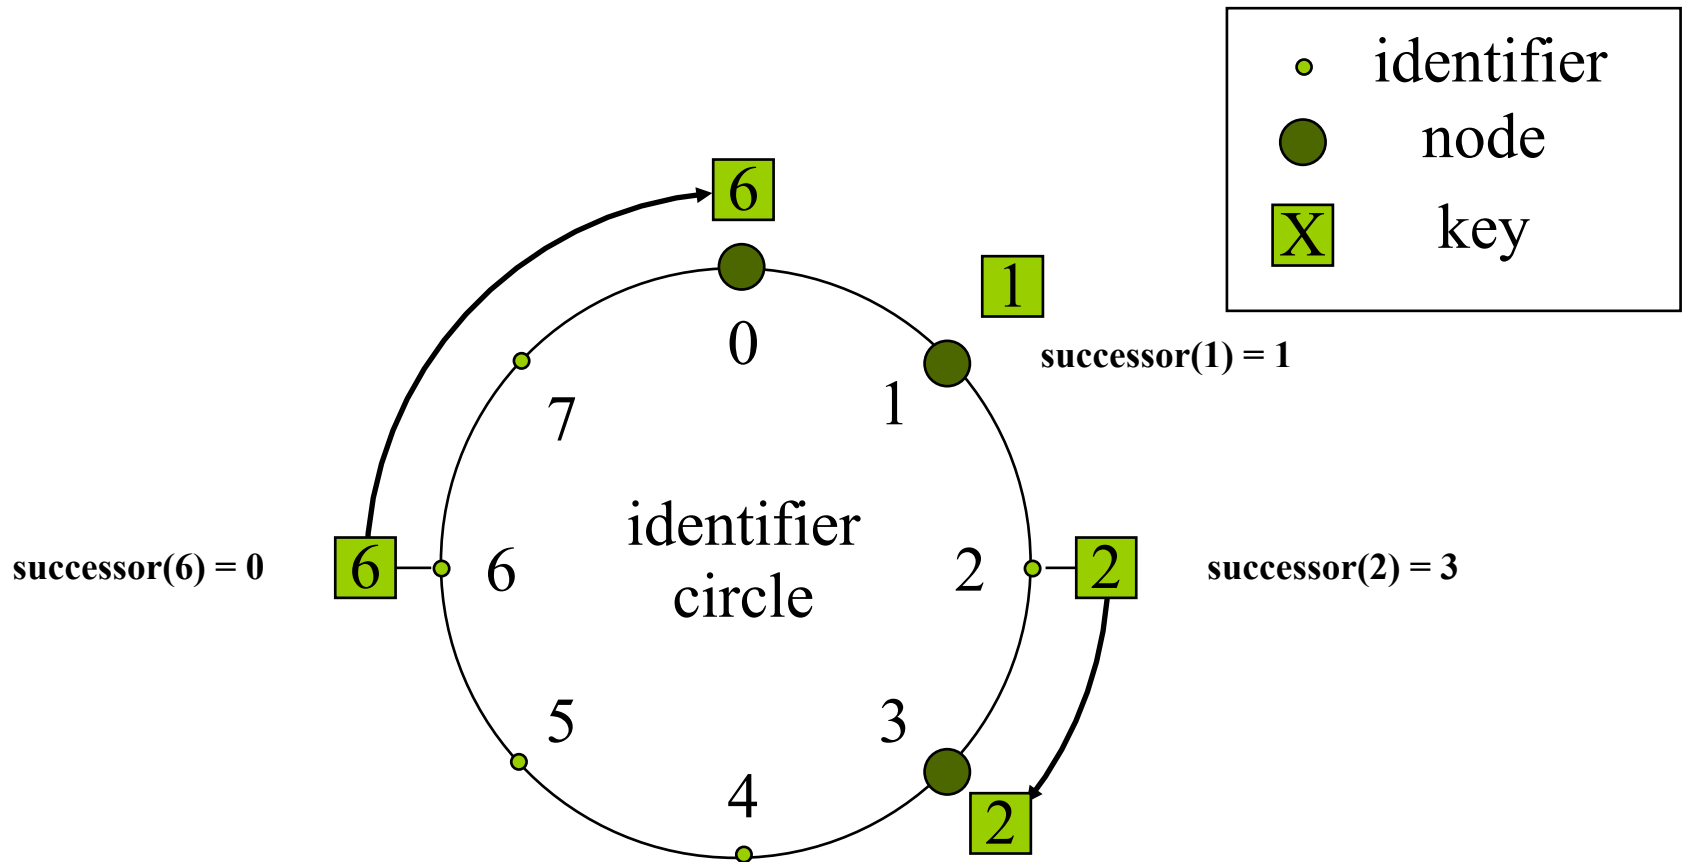


# DHT: Chord

---

- ▶ Consistent hashing is used as a base, with its properties

- ▶ Consistent hashing is used as a base, with its properties





# DHT: Chord

---

- ▶ In a distributed setting we need to route messages:
  - ▶ Need to store extra information to do so
    - ▶ A very small amount of routing information is necessary to implement consistent hashing
- ▶ Each node need only be aware of its successor node in the ring
- ▶ Queries for a given identifier can be passed around the circle via these pointers
- ▶ Is this efficient?





# DHT: Chord

---

- ▶ In a distributed setting we need to route messages:
  - ▶ Need to store extra information to do so
    - ▶ A very small amount of routing information is necessary to implement consistent hashing
- ▶ Each node need only be aware of its successor node in the ring
- ▶ Queries for a given identifier can be passed around the circle via these pointers
- ▶ What order of complexity does this have?



# DHT: Chord

---

- ▶ In a distributed setting we need to route messages:
  - ▶ Need to store extra information to do so
    - ▶ A very small amount of routing information is necessary to implement consistent hashing
- ▶ Each node need only be aware of its successor node in the ring
- ▶ Queries for a given identifier can be passed around the circle via these pointers
- ▶ What order of complexity does this have?
  - ▶ Resolution is correct but inefficient
    - ▶ May require traversing all nodes
      - $O(N)$

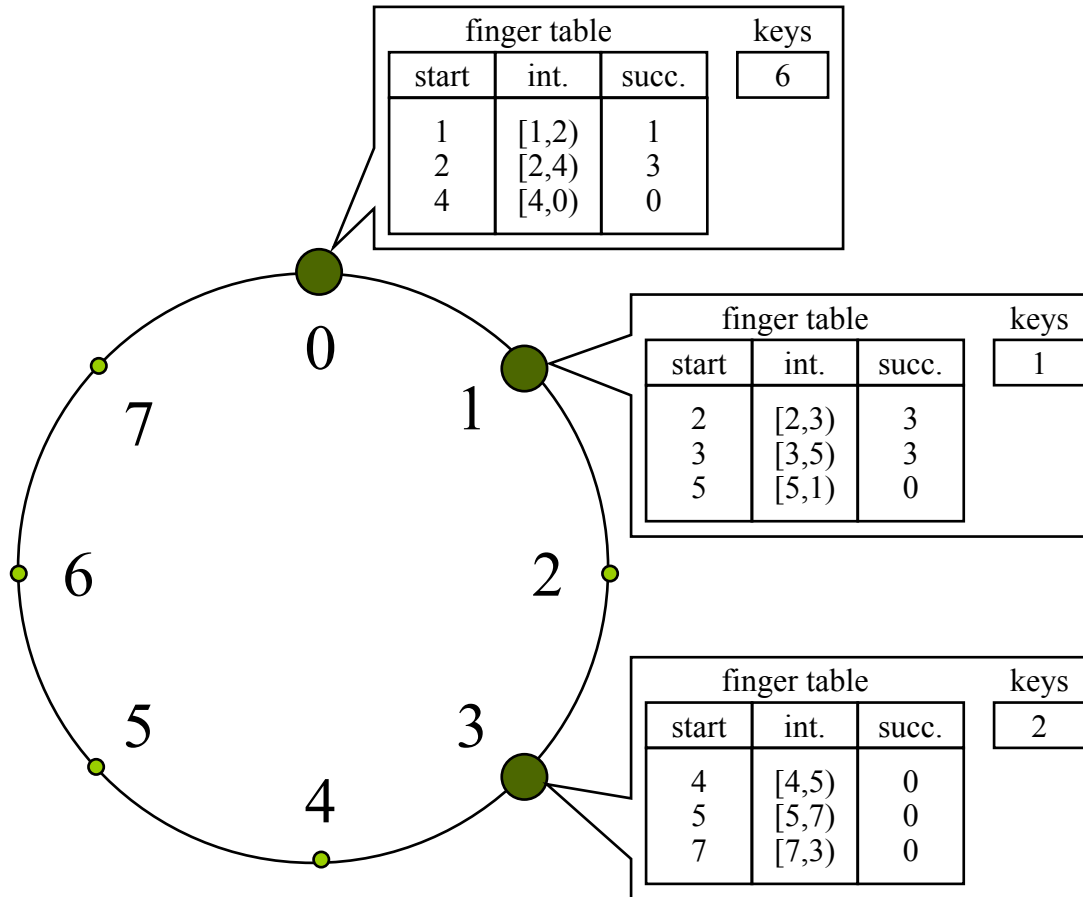


# DHT: Chord


---

- ▶ Will need extra information if we want to speedup the lookups
  - ▶ Didn't we say we could use some sort of a tree?...
  - ▶ How to do so in a distributed setting?


- ▶ Will need extra information if we want to speedup the lookups
  - ▶ Didn't we say we could use some sort of a tree?...
  - ▶ How to do so in a distributed setting?
- ▶ Each node maintains a routing table with  $\leq m$  entries
  - ▶ **Finger Table**
    - ▶  $N = 2^m$  or  $m = \log N$
- ▶ Small data structure even for large N
- ▶ For any node n, and entry i of its Finger Table, we have
  - ▶  $FT_i^n = \text{successor}(n + 2^{i-1}) = s$
  - ▶ S is called the  $i^{th}$  finger of node n.
    - ▶ `n.finger(i).node`



# DHT: Chord location




# DHT: Chord location







- ▶ Thus...
- ▶ Each node stores info about a small number of other nodes
  - ▶ Knows more about nearby nodes than about far away nodes
- ▶ Finger Table in a node has insufficient info to determine a successor of a key by itself
  - ▶ Repetitive queries are the key (no pun intended)
    - ▶ To nodes that immediately precede the given key
    - ▶ Eventually reach successor(key)












# DHT: Chord

---


- ▶ Basic “stabilization” protocol
  - ▶ Used to keep nodes’ successor pointers up to date
    - ▶ Sufficient to guarantee correctness of lookups
- ▶ Successor pointers can be used to verify finger table entries
- ▶ Every node runs *stabilize*
  - ▶ Periodically
  - ▶ To find new joined nodes


# Chord: Stabilization after Join




- **n joins**
  - predecessor = nil
  - n acquires  $n_s$  as successor via some  $n'$
  - n notifies  $n_s$  being the new predecessor
  - $n_s$  acquires n as its predecessor
- **$n_p$  runs stabilize**
  - $n_p$  asks  $n_s$  for its predecessor (now n)
  - $n_p$  acquires n as its successor
  - $n_p$  notifies n
  - n will acquire  $n_p$  as its predecessor
- **all predecessor and successor pointers are now correct**
- **fingers still need to be fixed, but old fingers will still work**

# Chord: Stabilization






- N26 joins the system
- N26 acquires N32 as its successor
- N26 notifies N32
- N32 acquires N26 as its predecessor



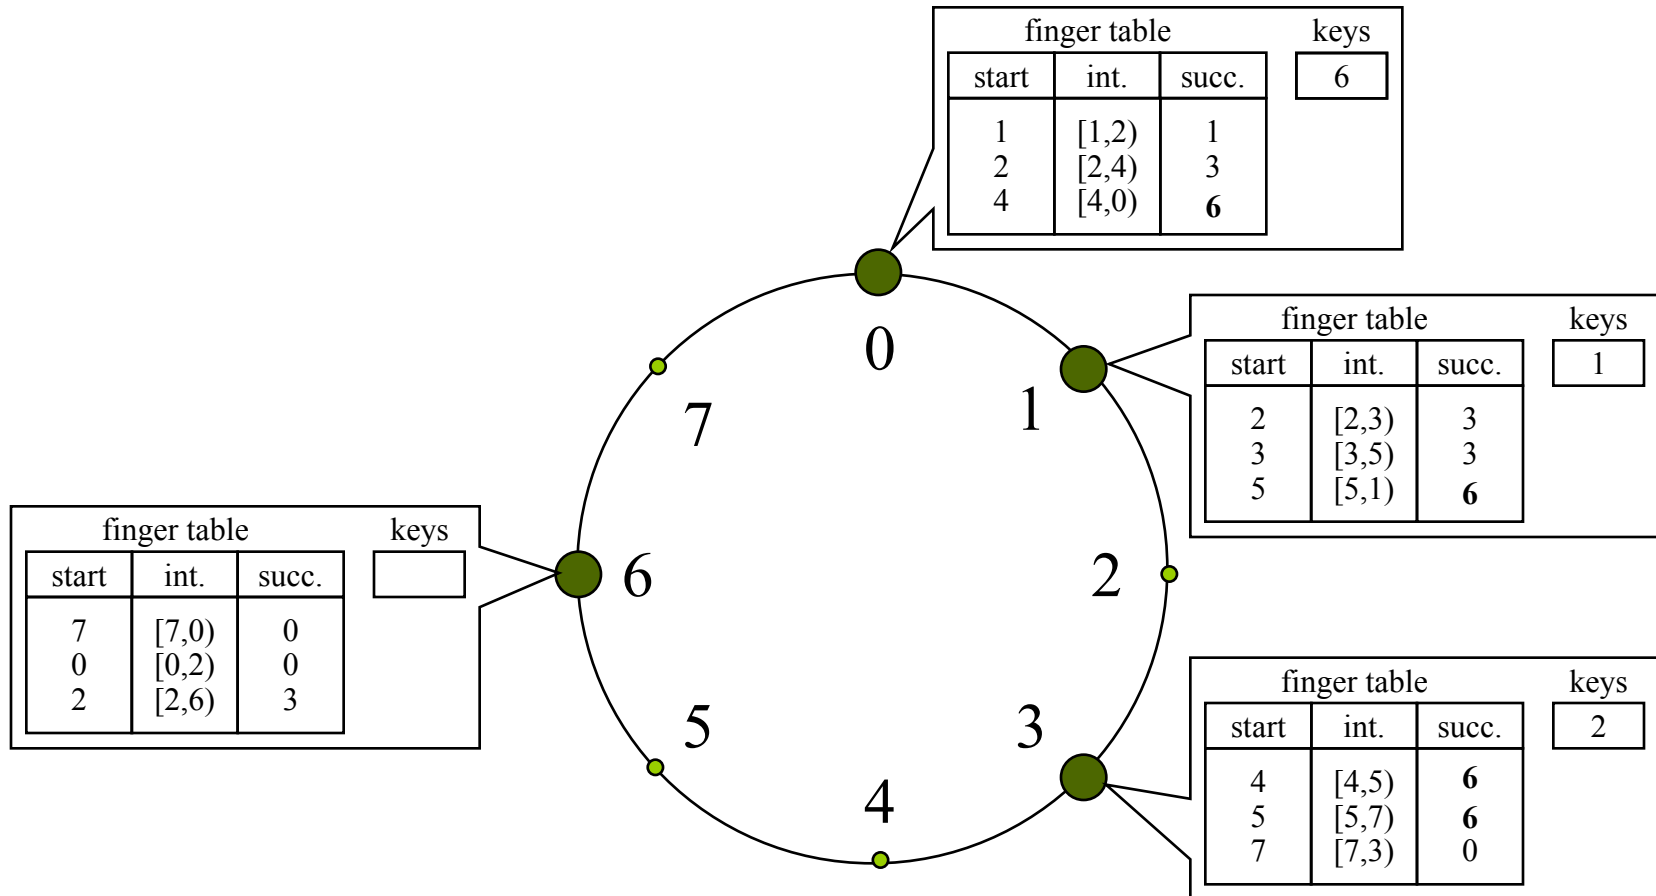
- N26 copies keys
- N21 runs stabilize() and asks its successor N32 for its predecessor which is N26.






- N21 acquires N26 as its successor
- N21 notifies N26 of its existence
- N26 acquires N21 as predecessor


# Chord: Node joins





Finger table	
N8 + 1	N14
N8 + 2	N14
N8 + 4	N14
N8 + 8	N21
N8 + 16	N32
N8 + 32	N42

- ▶ Correctness relies on correct successor pointers
- ▶ What happens, if N14, N21, N32 fail simultaneously?
- ▶ How can N8 acquire N38 as successor?



- ▶ Correctness relies on correct successor pointers
- ▶ What happens, if N14, N21, N32 fail simultaneously?
- ▶ How can N8 acquire N38 as successor?

Finger table

$N8 + 1$	N14
$N8 + 2$	N14
$N8 + 4$	N14
$N8 + 8$	N21
$N8 + 16$	N32
$N8 + 32$	N42

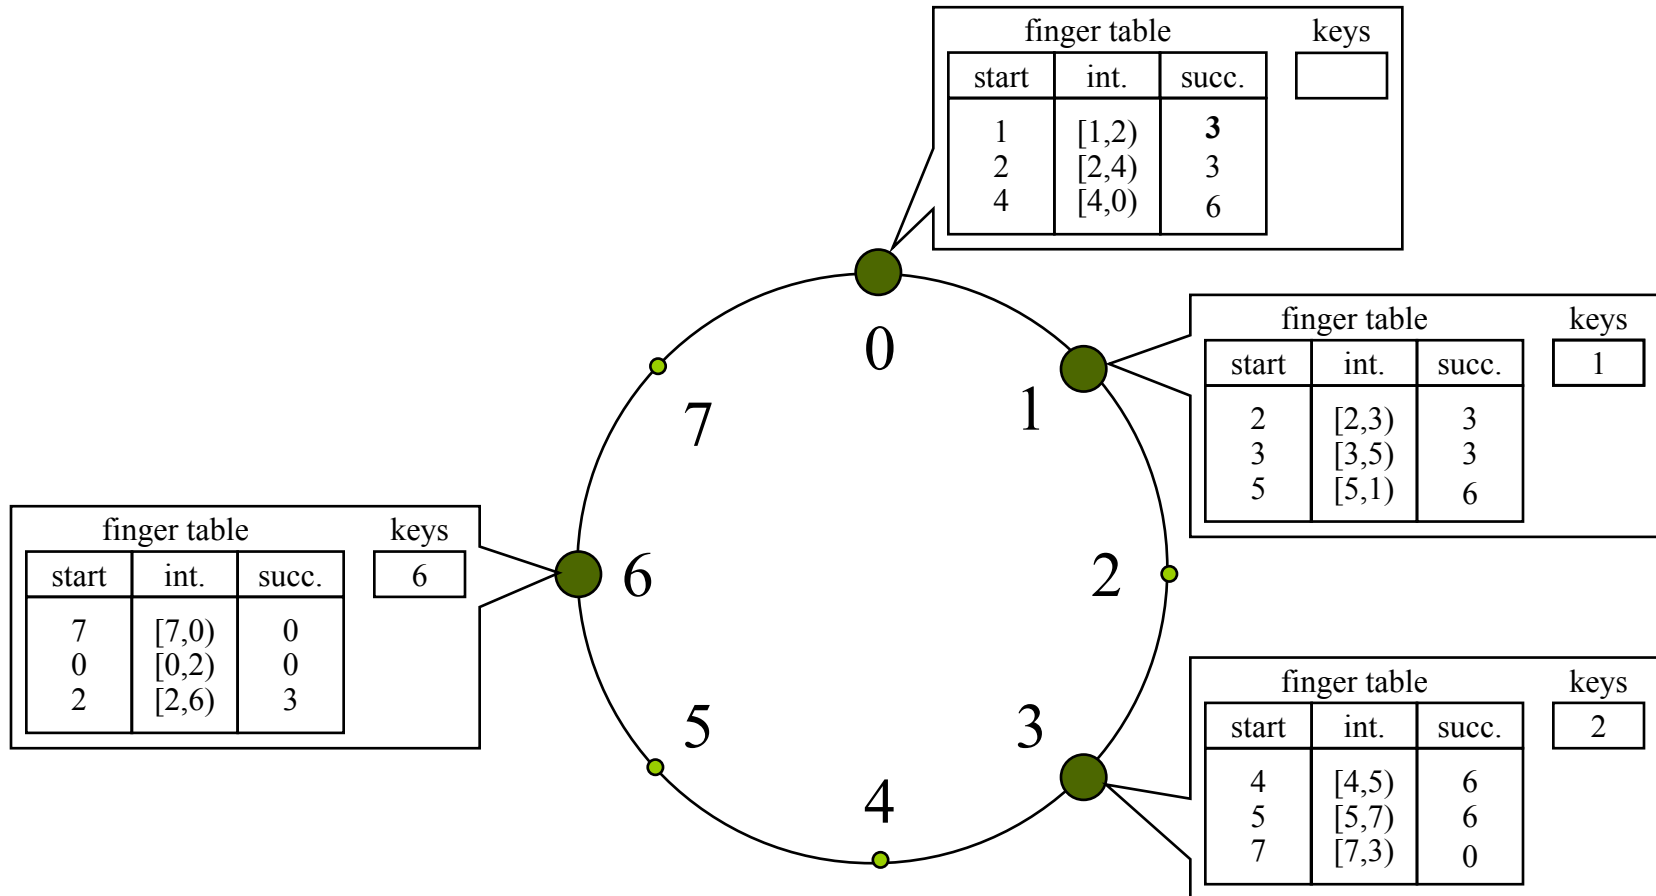


# DHT: Chord

---

- ▶ Failures. What do we do about them?
  - ▶ Key: Maintain correct successor pointers
- ▶ Each node maintains a *successor-list* of its  $r$  nearest successors on the ring
  - ▶ If node  $n$  notices its successor has failed...
    - ▶ It substitutes it with the first live entry in the list
  - ▶ *Stabilize* will correct finger table entries pointing to the failed node
    - ▶ Also *successor-list* entries
- ▶ Performance will depend on ratio of frequencies
  - ▶ Node joins/leaves
  - ▶ Invocation of stabilization protocol

# Chord: Node departures





# DHT: Chord

---

- ▶ Some numbers
  - ▶ As per Consistent Hashing
    - ▶ Each node takes care of  $K/N$  keys on average
    - ▶  $O(K/N)$  keys are relocated for joins/leaves
  - ▶ Lookups need  $O(\log N)$  messages
    - ▶ Thanks to the finger table
  - ▶ Reestablishing routing info and finger tables requires
    - ▶  $O(\log^2 N)$  messages