



Cloud Models



Motivation

- ▶ Main Models
 - ▶ IaaS (basic)
 - ▶ SaaS (motivation)
 - ▶ PaaS (utility)
 - ▶ Serverless/FaaS



- ▶ Foundational
 - ▶ Makes *utility computing* possible
 - ▶ Pay-as-you-go
 - Just like power/water/gas: utilities
- ▶ Fully depends on virtualization techniques
 - ▶ Versatility in resource management
- ▶ Fully depends on providers
 - ▶ On premise
 - ▶ Own CPD
 - ▶ Tailored CPD
 - ▶ Off-premise
 - ▶ Hybrid schemes
 - ▶ On-premise-then-off-premise



- ▶ SaaS KEEPS STATE through its lifetime
 - ▶ This is a service
 - ▶ The state is the history of event executions
 - ▶ It is a sort of DataBase
 - ▶ Each event is like a transaction that atomically modifies the state of the service.
- ▶ Thus
 - ▶ SaaS requires 24/7 activation
 - ▶ Supposedly is always active
 - ▶ Can be relaxed
 - ▶ And complex scaling logic
 - ▶ To meet QoS parameters
 - ▶ Same request at different times...
 - ▶ ...may produce different results



SaaS: is it always needed?

- ▶ Are there loads that do not require to store state?
 - ▶ I.e., purely functional
 - ▶ Fire and forget
- ▶ If so, do they require 24/7 activation?
 - ▶ On demand only
 - ▶ Called like a **F**unction (**as a S**ervice, of course)
 - No scaling issues, per se
 - Scaling is automatic
 - No state
 - Indeed we DO NOT WANT ANY STATE
 - ▶ Danger to break isolation
 - But possible latency issues when activating function




Serverless

- ▶ Aka FaaS: Function as a Service
 - ▶ Logic are small pieces of code
 - ▶ Respond to events
 - ▶ Event-driven
 - ▶ Functional
 - ▶ Logical activation when an even occurs
 - E.g. availability of a message on a message queue
 - Variety of triggers
- ▶ No local state store
 - ▶ But may use external persistent stores
 - ▶ i.e. SaaS (which keep state)
- ▶ Main advantage: ELASTIC BY DESIGN
 - ▶ If latencies are manageable

Load Example: ETL

- ▶ Extract, Transform, and Load
 - ▶ Basic process in Data Analysis tasks
 - ▶ When data arrives, it is conditioned/preprocessed
 - ▶ Then saved to a database


ETL without FaaS



Load Example: ETL, with SaaS

- ▶ Allocate an instance for 24/7 operation
 - ▶ Needed for when data arrives
 - ▶ Constant cost
 - ▶ Need to scale if lots of data arriving
- ▶ What happens if Data arrival is irregular?
 - ▶ Instance idle often
 - ▶ If suddenly lots of data arrive, potential scaling issues


ETL without FaaS



Load Example: ETL, with FaaS


- ▶ What happens if data arrives at irregular intervals?
 - ▶ Functions are activated only when there is data
 - ▶ Costs only when there is something to actually do
- ▶ If lots of data arrive suddenly
 - ▶ Multiple activations can be carried at the same time
 - ▶ Scaling is a non-issue for the function implementation
 - ▶ It impacts the FaaS implementation
 - ▶ It all depends on the framework

ETL with FaaS



Serverless Programming and SPAs

- ▶ Much of the application function carried out within an SPA
 - ▶ Static code downloaded from a site
 - ▶ E.g., S3/Github/Gitlab,...
- ▶ Interact with a back-end to perform some function
 - ▶ On demand, at irregular times
 - ▶ State maintained possible on a database of some sort (if at all)





Serverless

- ▶ Requires a framework
 - ▶ Configured to look for relevant events
 - ▶ And activate the code when an event arrives
 - ▶ Often provides a runtime for specific languages
 - ▶ Exceptions exist, and the pattern itself does not require it
- ▶ The framework itself must be stateful
 - ▶ Store the evnts
 - ▶ Keep track of the activations
 - ▶ Keep track of on who's behalf an activation is proceeding
 - ▶ Keep track of where to find the functions
- ▶ Framework is a SaaS
- ▶ Backend store is often necessary
 - ▶ Another SaaS



PaaS: Platform as a Service

- ▶ **Personas:**
 - ▶ **Developers**
 - ▶ **Integrators**
 - ▶ **Service managers**
- ▶ **Make life easy for all of them**
 - ▶ **Developers**
 - ▶ Focus on the app, forget systems-related noise
 - ▶ **Integrators**
 - ▶ Use a high level spec to convey structure of the microservices
 - ▶ **Service managers**
 - ▶ Let the platform automate most of the life cycle management



PaaS

▶ Automations:

- ▶ Fault-healing/high availability
- ▶ Scaling (horizontal/vertical)
- ▶ Disaster recovery
- ▶ Upgrade paths
- ▶ Security

▶ Approach

- ▶ Constrain how service applications are expressed
 - ▶ Give the platform extra knowledge to automate



▶ Benefits:

- ▶ Developers do not have to think about system
- ▶ Underlying OS is a non-issue
- ▶ Automations to adapt to changes in demand volume/types

▶ Drawbacks

- ▶ No control on underlying technology
- ▶ Must fit the framework
 - ▶ Opinionated way of structuring things...
 - ▶ Specific API for managing aspects of the app
- ▶ May be locked-in in some cases
 - ▶ Open source alternatives diminish this risk