



Virtualization



What is virtualization?

- ▶ Mechanism to substitute a hardware environment
 - ▶ For another, emulated environment
 - ▶ By any means
- ▶ Key to modern Cloud environments
 - ▶ Base for automation



Drivers for virtualization

- ▶ Expanding hardware capabilities
 - ▶ Allowing each single machine to do more simultaneous work
 - ▶ Need to increment utilization
- ▶ Server consolidation
 - ▶ Control costs
 - ▶ Simplify management
- ▶ Need to control large cluster installations
 - ▶ Server and render farms
- ▶ Improving security, reliability, and device independence
- ▶ Ability to run complex, OS-dependent applications in different hardware or OS environments



Goals

- ▶ Easily share a pool of computing resources
 - ▶ Elimination of manual resource management in CPDs
 - ▶ Increase reliability
 - ▶ Increase utilization of those resources
- ▶ Use same hardware for simultaneous different OS
 - ▶ E.g., run Linux on-windows (vice-versa), or on OSX, ...
 - ▶ Increase usefulness of computing resource
- ▶ Use same hardware to emulate even other hardware
 - ▶ Run ARM on i86 (or viceversa)



Isolation

- ▶ A zone's actions cannot influence another zone's actions
 - ▶ Derived from the fact that they share computing resources
 - ▶ Even its failures are isolated
- ▶ Two aspects:
 - ▶ Access:
 - ▶ What permissions can one get to manipulate a resource
 - ▶ How much information can one isolated “zone” capture from another isolated zone
 - ▶ Quotas:
 - ▶ How much of a resource can one zone consume



Access Isolation

- ▶ A zone cannot access a resource on another zone
 - ▶ File
 - ▶ User
 - ▶ Process
 - ▶ Core
 - ▶ Memory contents at an address
 - ▶ Messages being transmitted/received
- ▶ Common aspects for access aspect:
 - ▶ Resource can be named
 - ▶ They are identifiable
 - ▶ They are named somehow



Quota Isolation

- ▶ A zone cannot use more of a resource that it is given
 - ▶ Disk space
 - ▶ Memory
 - ▶ CPU usage
 - ▶ IOPs
 - ▶ Network bandwidth
- ▶ Common aspects for quota aspect:
 - ▶ No individual resource is named
 - ▶ They are NOT identifiable
 - ▶ E.g. we do not care what GB of ram we are given
 - ▶ Only magnitudes are important



Isolation Relaxations

- ▶ Establish policies for accessing individual resources among zones
 - ▶ ACLs
 - ▶ Capabilities based on name resolutions
 - ▶ Problem:
 - ▶ Ensure the policies can be clearly established and properly enforced
- ▶ Permit limited competition for shared resources:
 - ▶ Map multiple pools of resources to the same underlying hardware resource



Virtualization approaches

- ▶ Two major types
 - ▶ Hardware/Machine virtualization
 - ▶ OS/Light virtualization
- ▶ Also other kinds of virtualization
 - ▶ Network Virtualization



Machine/Hardware virtualization

- ▶ The hardware is “emulated” by a piece of software
 - ▶ Originally, literally an interpreter of machine code
 - ▶ With enormous loss of performance
 - ▶ With high degree of access isolation
 - ▶ With poor degree of quota isolation
 - ▶ Main Goal:
 - Run different environments (Windows/Unix,...)
 - Run different HW architectures (ARM on Intel,...)
- ▶ Eventually allows direct access to the CPU
 - ▶ Still, access to other hardware parts limited
 - ▶ Software in charge: Hypervisor
 - ▶ Enhancement: “enlightened” kernel modules, calling Hypervisor
- ▶ Eventually, hardware support with specific machine instructions
 - ▶ Run by the Hypervisor/Host system.
- ▶ Eventually, hierarchical virtualization.



Machine/Hardware virtualization

- ▶ Produces what are known as VMs (Virtual Machines)
 - ▶ Access: individual Memory addresses
 - ▶ According to quotas
 - ▶ Individual devices
 - ▶ Individual disks and/or partitions
 - ▶ Individual cores
 - ▶ According to quotas
- ▶ Perfect partition/Overcommitting:
 - ▶ Maps each hardware resource to each VM
 - ▶ VM has exclusive access to hardware resource
 - ▶ More simulated resources than hardware resources (quota relaxation)
 - ▶ Different VMs compete for resources in a limited way



Machine/Hardware virtualization

- ▶ Two types of hypervisors: Type I and Type 2
- ▶ Type-I: Directly on Hardware
 - ▶ Example Hypervisors:
 - ▶ VMWARE ESXi
 - ▶ HYPER-V
 - ▶ Xen
 - ▶ IBM CP/CMS (oldie), and z/VM
 - ▶ Oracle/SPARC VM Server
 - ▶ ...
- ▶ Type-II: Hosted on OS
 - ▶ VirtualBox
 - ▶ Parallels
 - ▶ Qemu
 - ▶ WMware Desktop/Player
- ▶ Not always clear distinction: KVM (Linux kernel module)



Network virtualization

- ▶ At various layers
 - ▶ E.g., layer2 physical network is simulated
 - ▶ Various flavors of virtual switches
 - ▶ Example: various vswitches
 - ▶ Example: various veth interfaces
 - ▶ Example: zerotier-one
 - ▶ Also layer3
 - ▶ Example: Wireguard
- ▶ Typically uses encapsulation techniques
- ▶ Gives rise to soft networks



OS/Light virtualization

- ▶ Oldest technique for time sharing
 - ▶ Processes
 - ▶ Virtual memory
 - ▶ File systems
 - ▶ ...
- ▶ Isolation imperfections abound
 - ▶ Shared name spaces
 - ▶ Files/Process IDs/users/sockets
 - ▶ ACLs to limit access, but still wide potential holes in isolation
 - ▶ Shared software stack & runtimes
 - ▶ Libraries
 - ▶ Daemons
 - ▶ ...
 - ▶ Competition for resources
 - ▶ Memory/CPU/IO
 - ▶ Limited quota enforcement mechanisms



OS/Light virtualization

▶ Consequence:

- ▶ Large API surface to protect!!!
- ▶ Too much relaxation in access (naming)
- ▶ Too much relaxation in resource usage (quotas)

▶ Requirements

- ▶ Severely reduce API surface
- ▶ Fully isolate name spaces from which to obtain access to specific resource
- ▶ Effectively cap the amount of resources an environment can make use of.



OS/Light virtualization

- ▶ Severely reduces surface AREA:
 - ▶ Only the kernel is shared
 - ▶ The rest of the software stack is **NOT** shared.
 - ▶ Kernel calls are the only surface area exposed.
- ▶ Kernel imposed namespace separation
 - ▶ Kernel is shared and it governs the rest of the OS
 - ▶ Kernel imposes the lowest level naming of objects
 - ▶ Mount points for file systems
 - ▶ Users, Groups, Process IDs, Network, devices, sockets,...
- ▶ Kernel imposed Resource consumption control
 - ▶ Quotas for memory/CPU/IOPS/Bandwidth
 - ▶ Mechanism to impose them on environments
- ▶ Environments built out of groups of processes on the OS.




Machine vs Light virtualization

▶ VMs

- ▶ Need to boot a complete OS (the guest OS)
 - ▶ Boot process identical to that performed by a physical machine
 - ▶ Takes time
 - ▶ Takes more resources to simulate the machine environment
 - ▶ In general, needs to start many processes to provide the machine environment
- ▶ Host OS can be completely different from guest OS
 - ▶ E.g. Windows/Linux
 - Windows/MacOSX
 - MacOSX/Linux
 - MacOSX/Windows
 - Linux/*
- ▶ Code cannot be shared among different guests

Machine vs Light virtualization






Machine vs Light virtualization

- ▶ Container-like entities (light virtualization)
 - ▶ “Simply” needs to set up the isolation environment
 - ▶ Needs support from the OS kernel
 - Namespace isolation
 - Thus communication primitives isolation
 - ▶ Including IPC mechanisms
 - ▶ Including virtualized network interfaces
 - User space isolation
 - Superuser (root in Linux) inside has nothing to do with root at the host
 - Is a special case of namespace isolation
 - Resource capping
 - CPU
 - Memory
 - ...
 - ▶ Kernel code can be shared among containers
 - Other code can also be shared... at risk
 - ▶ Constraint: **guest must work on Host's kernel API**



Machine vs Light virtualization






Machine vs Light virtualization

- ▶ Approaches can be mixed
 - ▶ Light virtualization inside Machine virtualization

Machine vs Light virtualization






Machine vs Light virtualization

- ▶ Network virtualization
 - ▶ Similar in both cases
 - ▶ Uses the concept of a virtual switch at the host
 - ▶ The virtual switch is “connected” to virtual ethernet interfaces
 - ▶ In light virtualization, the OS kernel must support virtualized interfaces
- ▶ The virtual switch manages
 - ▶ Addressability
 - ▶ What IP networks are reachable
 - ▶ Available bandwidth: rate limiting
 - ▶ How many MB/s are available
 - ▶ In some fancier set-ups
 - ▶ Throtling

Machine vs Light virtualization





Containers: History

- ▶ FreeBSD Jails (~2000)
- ▶ Solaris zones (~2005)
 - ▶ Windows NT jails (unreleased) (2006)
- ▶ Linux mechanisms
 - ▶ 2006-2008
 - ▶ LXC libraries
 - ▶ Docker



Light virtualization on Linux: containers

- ▶ Supported by the kernel
 - ▶ Kernel version ≥ 3.6
- ▶ Based on kernel modules for
 - ▶ Namespace isolation
 - ▶ namespaces
 - ▶ Resource quota handling
 - ▶ cgroups



Linux cgroups

- ▶ cgroups == **Control Groups**
- ▶ Kernel mechanism to limit resource consumption
 - ▶ By measuring consumption and configuring quotas
 - ▶ Applied to a group of processes
 - ▶ All bound to the same limits
 - ▶ Can be further restricted in the process tree
 - Hierarchical
 - Limits set for a group are inherited by subgroups
- ▶ Generic mechanism
 - ▶ Can be applied to any process in the system
 - ▶ Nothing special about “containers”
 - ▶ Very granular



Linux cgroups

- ▶ Provide the following functionalities
- ▶ **Resource limiting (as mentioned)**
 - ▶ A group can be configured not to exceed a specified memory limit or use more than the desired amount of processors or be limited to specific peripheral devices.
- ▶ **Prioritization:**
 - ▶ One or more groups may be configured to utilize fewer or more CPUs or disk I/O throughput.
- ▶ **Accounting:**
 - ▶ A group's resource usage is monitored and measured.
- ▶ **Control:**
 - ▶ Groups of processes can be frozen or stopped and restarted.



Linux cgroups

- ▶ Cgroup subsystems (controllers)
- ▶ Each manages a kind of resource
 - ▶ Responsible for distributing it to the processes
 - ▶ E.g. *memory* controller restricts memory usage
 - ▶ E.g. *cpuacct* controller restricts cpu usage
- ▶ Controlled through *cgroup* filesystem
 - ▶ Mounted at `/sys/fs/cgroup`
 - ▶ Controllers mounted within it
 - ▶ Each group is a subdir of a controller mount point



Linux cgroups: example memory

- ▶ Create a cgroup example
 - ▶ `/sys/fs/cgroup/memory/example`
 - ▶ To limit it, we must add a file
 - ▶ `$ echo 50000000 | sudo tee /sys/fs/cgroup/memory/foo/memory.limit_in_bytes`
- ▶ Consider the app:
 - ▶ `$ cat test.sh`
 - ▶ `#!/bin/sh`
 - ▶ `while [1]; do echo "hello world" sleep 60 done`
 - ▶ Run it in the background
 - ▶ `$ sh ~/test.sh &`
 - ▶ Find its PID (e.g. 1234) and place it within the cgroup
 - ▶ `Echo 1234 > /sys/fs/cgroup/memory/example/cgroup.procs`
 - ▶ Try:
 - ▶ `$ ps -o cgroup 1448`



Linux namespaces

- ▶ Mechanism to limit visibility
 - ▶ Applied to a group of processes
- ▶ Types of resources covered:
 - ▶ Process trees (their ids)
 - ▶ Network interfaces
 - ▶ User Ids/Group Ids
 - ▶ File System mount points
 - ▶ Includes Unix Sockets
 - ▶ Devices
- ▶ Generic mechanism
 - ▶ Can be applied to any process in a system
 - ▶ Granular



Linux namespaces

▶ lsns

▶ Command lists the set of available namespaces on a system

```
▶ 4026531835 cgroup      5    938 ubuntu /lib/systemd/systemd --user
▶ 4026531836 pid        5    938 ubuntu /lib/systemd/systemd --user
▶ 4026531837 user       5    938 ubuntu /lib/systemd/systemd --user
▶ 4026531838 uts        5    938 ubuntu /lib/systemd/systemd --user
▶ 4026531839 ipc        5    938 ubuntu /lib/systemd/systemd --user
▶ 4026531840 mnt         5    938 ubuntu /lib/systemd/systemd --user
▶ 4026531992 net         5    938 ubuntu /lib/systemd/systemd --user
```

▶ Each process has a PID within a PID namespace

▶ Processes in different PID namespaces cannot interact with each other by PID

- ▶ They make no sense across PID namespaces

- ▶ Similar to the namespace of file descriptors

- ☐ Each process has its own, and cannot transfer them to another process



Linux namespaces

▶ The unshare command

- ▶ It runs a subprocess with the given executable in their own namespace
 - ▶ It indicates the namespaces that ran process MUST NOT share with the rest of the system

▶ Example

- ▶ `sudo unshare --fork --pid --mount-proc zsh`
- ▶ Runs a zsh subprocess, and offers a shell prompt
 - ▶ Any process launched from that prompt will be within the newly created namespace for zsh
 - ▶ Try: `pidof zsh`
 - You get `I`
 - ▶ On another terminal, **pidof zsh**, the result is different
- ▶ Many options available
 - ▶ E.g., mapping root /proc, ...

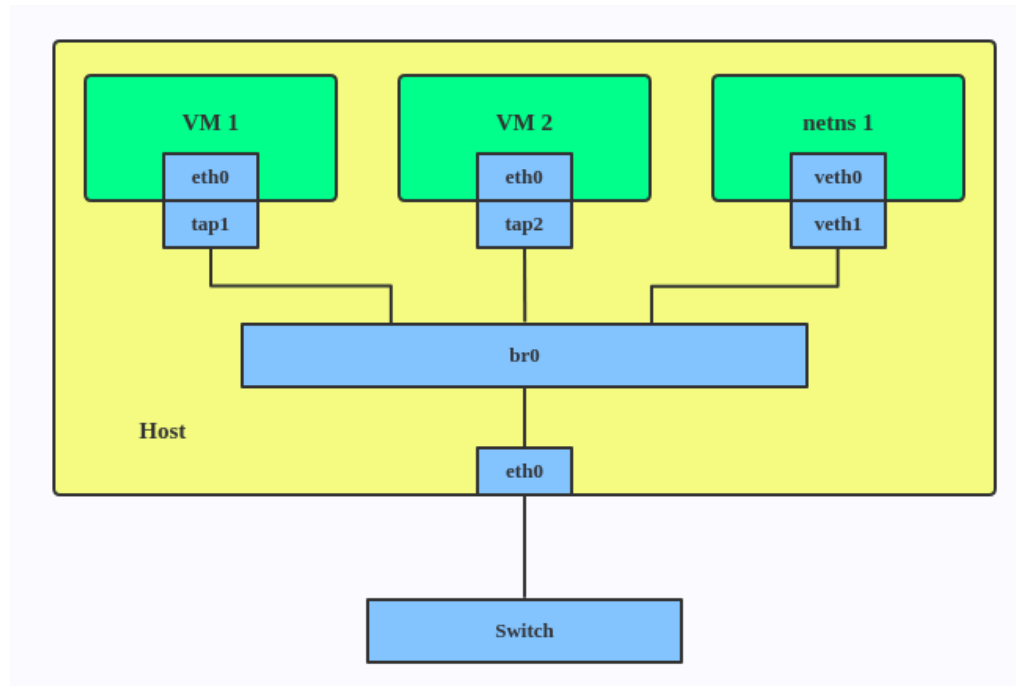


Virtual network interfaces

- ▶ Special virtualized network devices
- ▶ Interact with network namespace
- ▶ Many kinds in the Linux kernel (use `ip link help`)
 - Bridge
 - Bonded interface
 - Team device
 - VLAN (Virtual LAN)
 - VXLAN (Virtual eXtensible Local Area Network)
 - MACVLAN
 - IPVLAN
 - MACVTAP/IPVTAP
 - MACsec (Media Access Control Security)
 - VETH (Virtual Ethernet)
 - VCAN (Virtual CAN)
 - VXCAN (Virtual CAN tunnel)
 - IPOIB (IP-over-InfiniBand)
 - NLMON (NetLink MONitor)
 - Dummy interface
 - IFB (Intermediate Functional Block)
 - netdevsim

Network virtualization: Bridge

- ▶ Behaves like a network Switch
- ▶ Forwards packets among connected interfaces
- ▶ Use a bridge when you want to establish communication channels between VMs, containers, and your hosts.



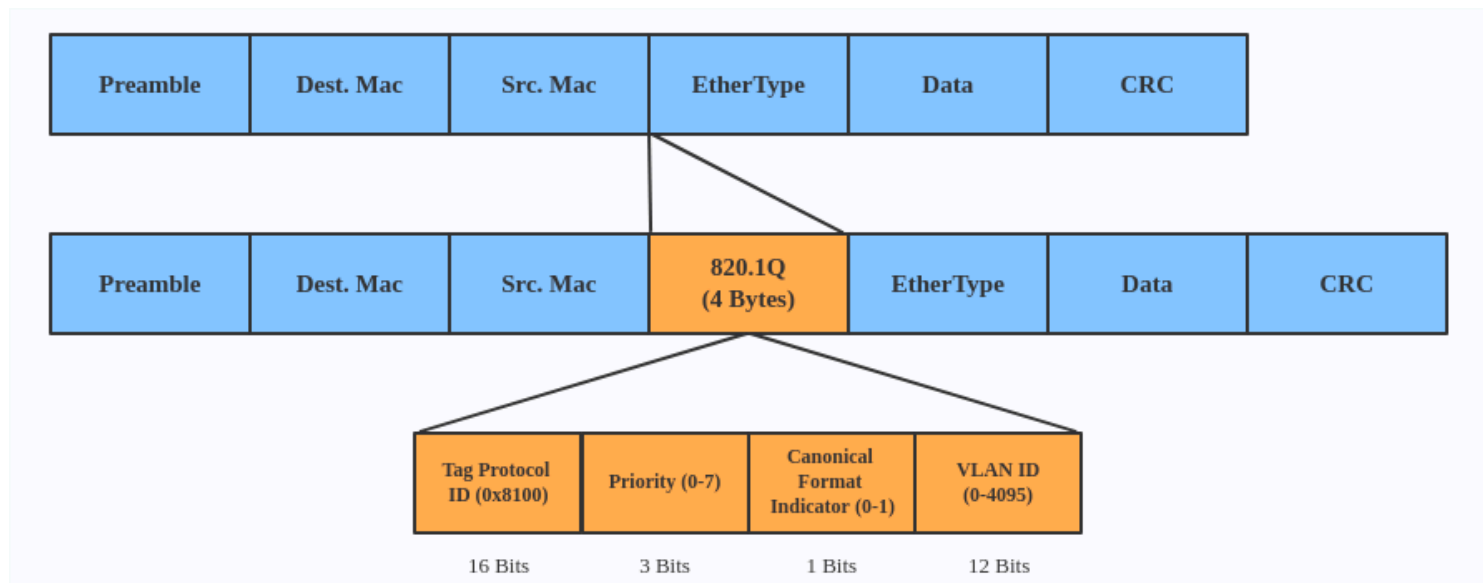


Network virtualization: Bridge

```
# ip link add br0 type bridge  
# ip link set eth0 master br0  
# ip link set tap1 master br0  
# ip link set tap2 master br0  
# ip link set veth1 master br0
```


Network virtualization: VLAN

- ▶ “Virtual LAN”
- ▶ Separates L2 (e.g. ethernet) domains by adding tags to network packets
- ▶ Use same physical hardware to form multiple L2 domains
- ▶ Look as different ethernet interfaces



Network virtualization: VLAN

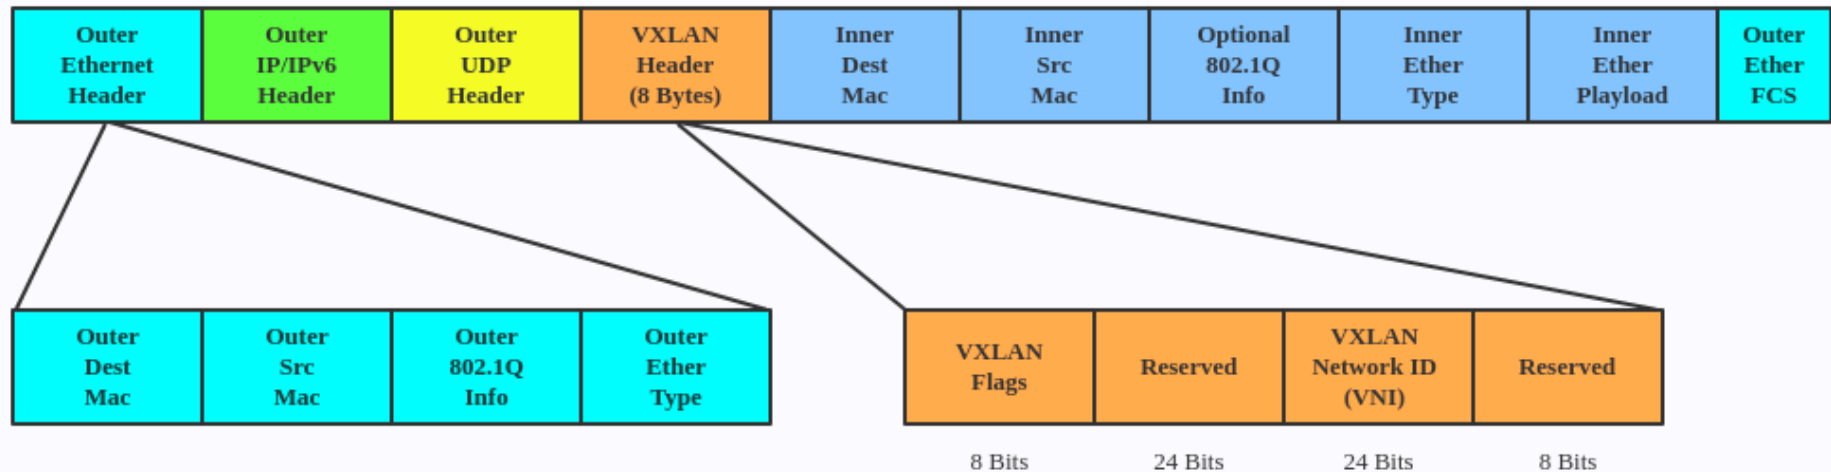
```
# ip link add link eth0 name eth0.2 type vlan id 2  
# ip link add link eth0 name eth0.3 type vlan id 3
```



- Look as different ethernet interfaces from within


Network virtualization: VXLAN

- ▶ “Virtual eXtensible LAN”
- ▶ **TUNNELING** protocol
 - ▶ Solves limited VLAN ids
 - ▶ VXLAN ID (segment ID) is 24 bits long → 16,777,216 virtual LANs
 - ▶ 4,096 times the VLAN capacity.
- ▶ Encapsulates Layer 2 frames with a VXLAN header into a UDP-IP packet



Network virtualization: VXLAN

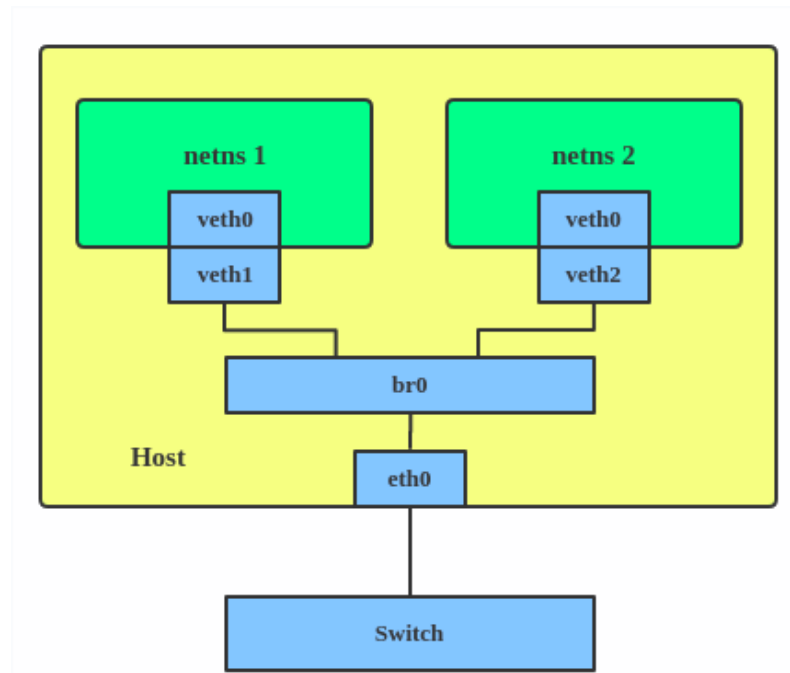
```
# ip link add vx0 type vxlan id 100 local 1.1.1.1 remote 2.2.2.2 dev eth0 dstport 4789
```



- ▶ Typically deployed in data centers on virtualized hosts
 - ▶ May be spread across multiple racks.


Network virtualization: MACVLAN

- ▶ Multiple interfaces on top of a physical one
 - ▶ Each has a different MAC address
- ▶ Without MACVLAN you need bridges to access the network
 - ▶ Plus, other virtualized interfaces (e.g. Veth)



Network virtualization: MACVLAN


- ▶ With MACVLAN
- ▶ Directly attach the MACVLAN interface to the namespace where it should be used



Network virtualization: Types of MACVLAN

► Private


- Forbids communication among interfaces on same physical interface



Network virtualization: Types of MACVLAN

► VEPA


- Communication through the external switch
 - Must support *hairpin* mode
 - Or... There is a router that redirects packets back



Network virtualization: Types of MACVLAN

► BRIDGE


- Phys interface behaves as a simple bridge also for internal macvlans
- Packets do not have to travel outside



Network virtualization: Types of MACVLAN

► PassThrough

- Allows a single VM/namespace to be connected directly to the physical interface





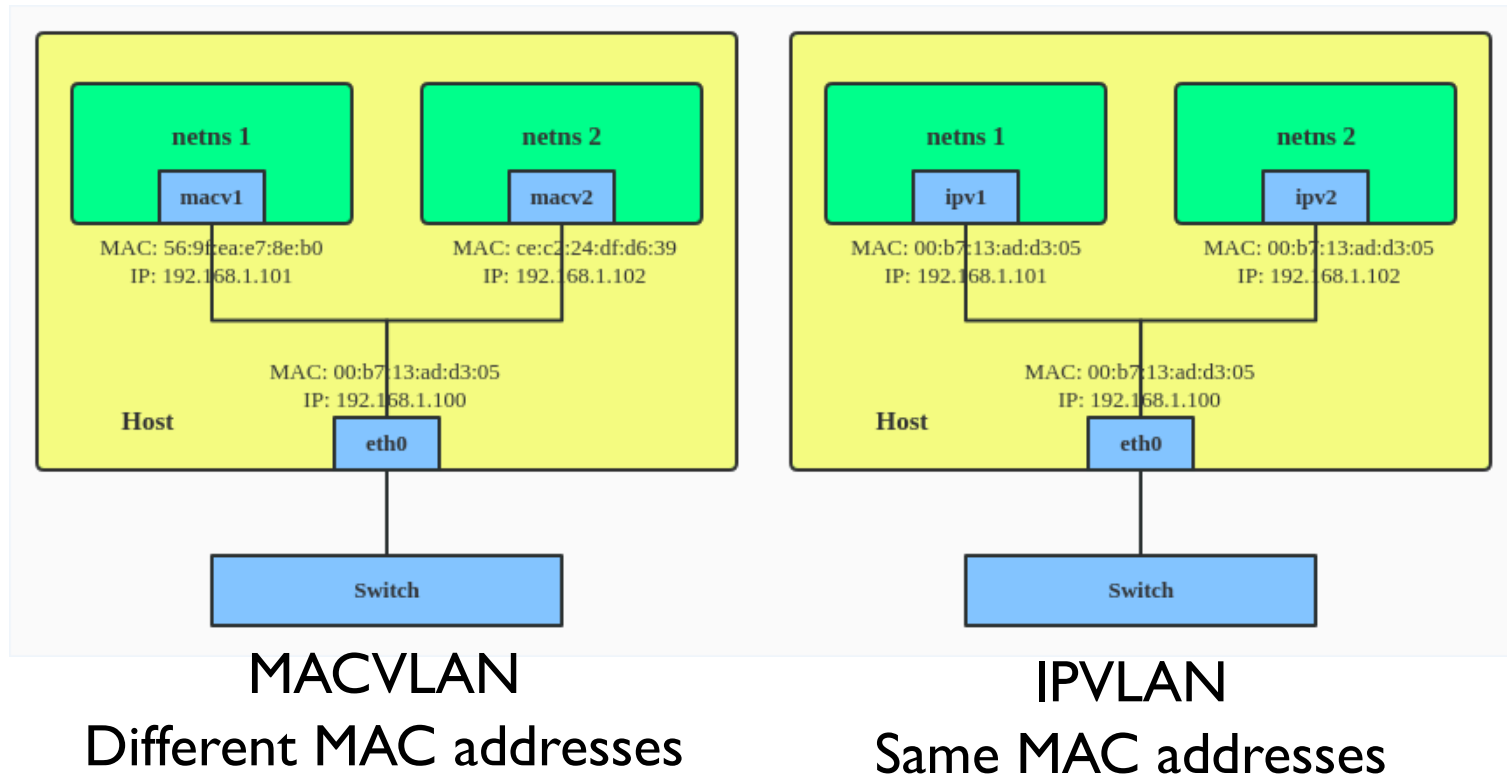
Network virtualization: Summary MACVLAN

- ▶ Use MACVLAN when connecting directly to a physical network from containers
- ▶ BRIDGE is the most common mode

```
# ip link add macvlan1 link eth0 type macvlan mode bridge
# ip link add macvlan2 link eth0 type macvlan mode bridge
# ip netns add net1
# ip netns add net2
# ip link set macvlan1 netns net1
# ip link set macvlan2 netns net2
```


Network virtualization: IPVLAN

- ▶ Similar to MACVLAN
 - ▶ But Interfaces share MAC address




Network virtualization: IPVLAN

- ▶ Two modes: Mode L2 or Mode L3
- ▶ In mode L2, it behaves like MACVLAN BRIDGE
 - ▶ Parent interface behaves like a simple bridge



Network virtualization: IPVLAN


- ▶ In mode L3,
 - ▶ Parent interface behaves like a router
 - ▶ Potentially better scalability





Network virtualization: IPVLAN

```
# ip netns add netns1
# ip link add ipv1 link eth0 type ipvlan mode l2
# ip link set dev ipv1 netns netns1
# ip netns add netns2
# ip link add ipv2 link eth0 type ipvlan mode l2
# ip link set dev ipv2 netns netns2
```





Network virtualization: MACVLAN or IPVLAN

- ▶ When to select IPVLAN over MACVLAN:
 - ▶ The host connected to the external switch/router has policy configured that allows only one mac per port.
 - ▶ Number of virtual devices created on a master exceed the MAC capacity
 - ▶ Forces NIC in promiscuous mode
 - Impact on performance
 - ▶ The slave device is to be put into an untrusted namespace
 - ▶ L2 on the slave could be misused.




Network virtualization: VETH

- ▶ Virtual EThernet: local ethernet tunnel
 - ▶ No tunnel traffic exits the host
- ▶ Created as pairs
 - ▶ Traffic entering one device immediately exits the other device
 - ▶ Like a cable
- ▶ When either device is down, the whole link is down
- ▶ Use VETH when network namespaces need to communicate
 - ▶ With the host
 - ▶ Among themselves



Network virtualization: VETH

```
# ip netns add netns1
# ip netns add netns2
# ip link add veth1 netns netns1 type veth peer name veth1
netns netns2
# ip link add veth0 netns netns1 type veth peer name veth1
# ip link add veth0 netns netns2 type veth peer name veth2
```





Network virtualization: DUMMY

- ▶ Fully virtual
 - ▶ E.g. The loopback interface is also fully virtual
- ▶ For routing packets through
 - ▶ With no transmission
- ▶ Used for testing/debugging

```
# ip link add dummy1 type dummy
# ip addr add 192.168.1.1/24 dev dummy1
# ip link set dummy1 up
```



Linux Containers

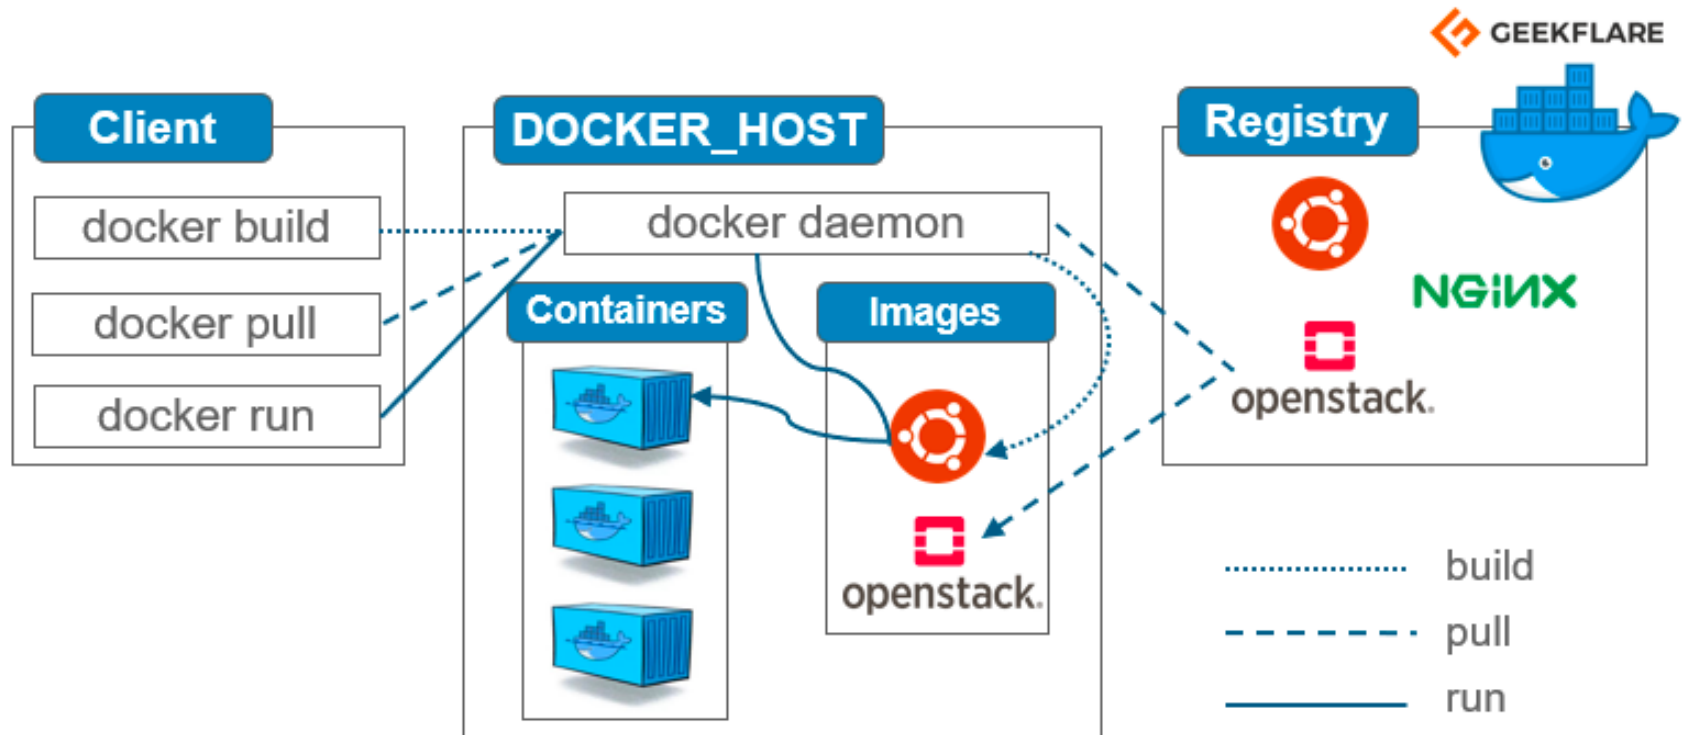
- ▶ Group of processes subjected to
 - ▶ Full set of cgroup policies
 - ▶ All resources are contemplated
 - ▶ Full set of namespace policies
 - ▶ All namespaces are contemplated
- ▶ Full cgroup/namespace policies isolate a container
 - ▶ Specifically from other containers
 - ▶ Two-way isolation
 - ▶ One-way isolation of the host
 - ▶ Containers cannot access the host but host can access containers
 - Host is privileged wrt kernel: full access to kernel API
- ▶ Various isolation relaxations
 - ▶ E.g. Share named resources
 - ▶ E.g. non-quotaed resources



Docker

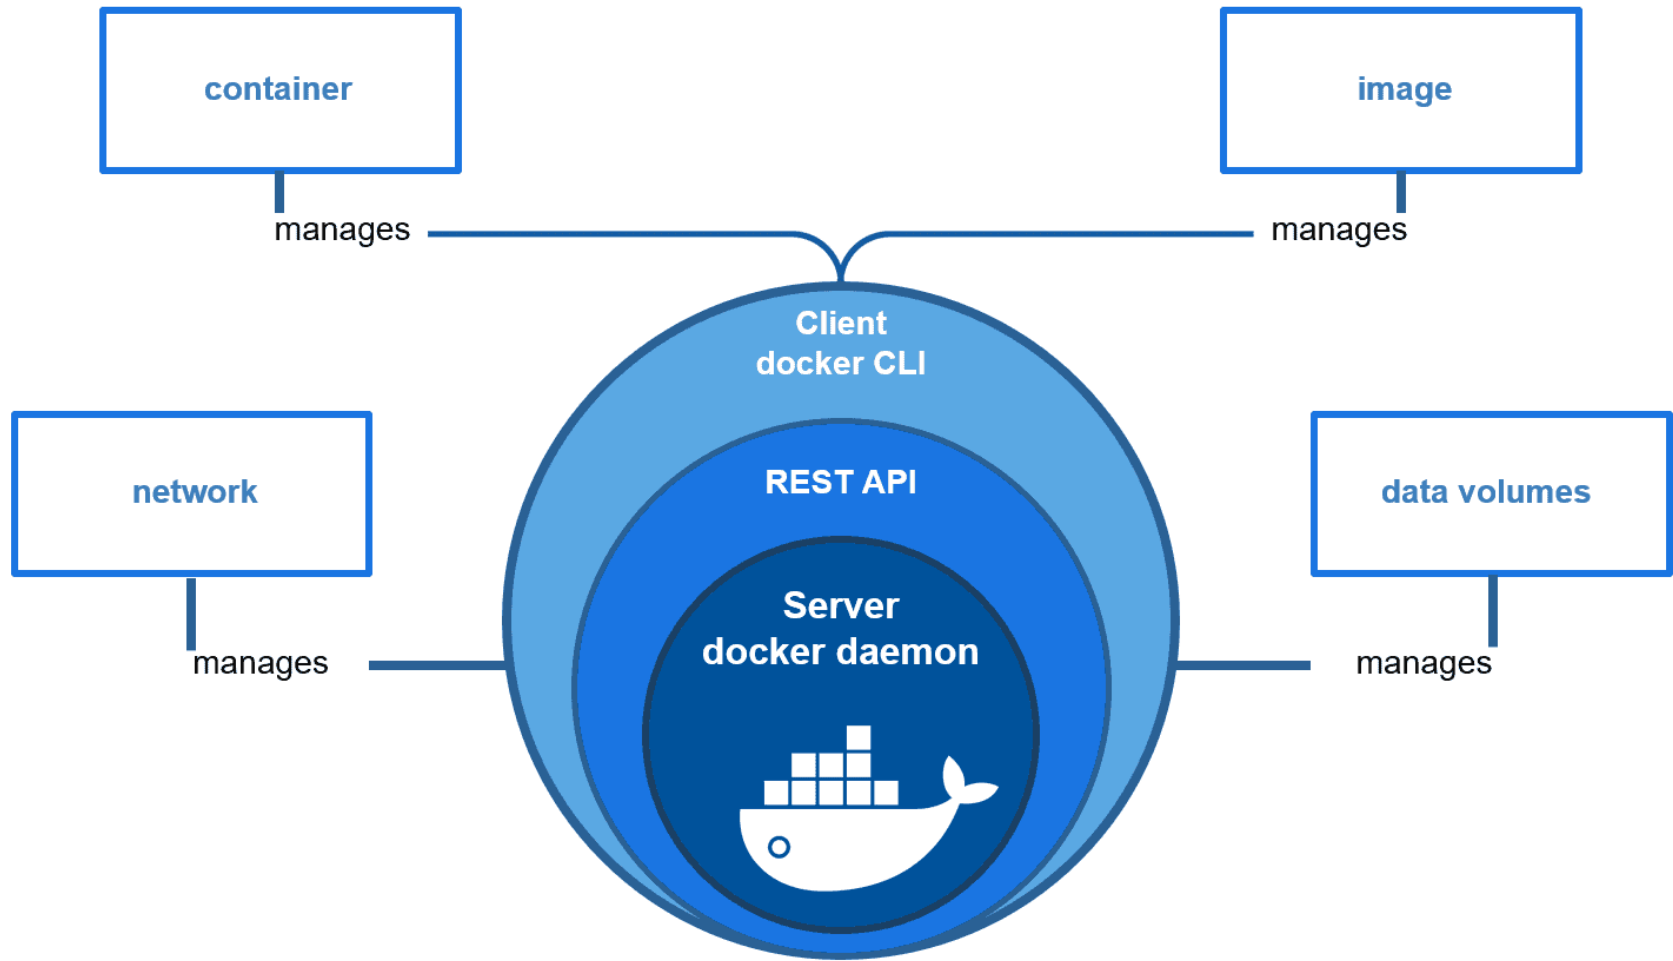
- ▶ A container management system
- ▶ Flexible architecture
 - ▶ Daemon (dockerd)
 - ▶ CLI command (docker)
 - ▶ Various plug-ins
 - ▶ Storage
 - ▶ Networking
 - ▶ An image standard
 - ▶ Based on a layered file system
 - ▶ Leveraged to port software stacks among hosts
 - ▶ Caches images on local host

Docker High level architecture





Docker Conceptual structure





Docker warm up exercises

► Go to

<https://github.com/eficode-academy/docker-katas/tree/master/labs>