

Máster Universitario en Ingeniería Informática

Sistemas Inteligentes

Unit 5. N-grams and FSA - Theory

2022/2023



1. N-grams
2. Introduction to Probabilistic Finite-State Automata
3. Relation between N-grams and FSA
4. References



1. N-grams

2. Introduction to Probabilistic Finite-State Automata

3. Relation between N-grams and FSA

4. References



- ▶ Assign a probability score to every word sequence
- ▶ Reflect a previous knowledge of a text source, predicting the most likely occurrence of words using its context

The observation probability of a word sequence $\mathbf{w} = \langle w_1 \dots w_m \rangle$:

$$P(\mathbf{w}) = P(w_1) \cdot \prod_{i=2}^m P(w_i | w_1 \dots w_{i-1}) \quad \text{with } w_i \in \Sigma$$

where $P(w_i | w_1 \dots w_{i-1})$ is the probability of having a word w_i , given a previous word history $w_1 \dots w_{i-1}$

The estimation of $P(\mathbf{w})$ is prohibitively expensive when vocabulary $|\Sigma|$ size becomes huge: $|\Sigma|^{i-1}$ different possible histories



- ▶ Assign a probability score to every word sequence
- ▶ Reflect a previous knowledge of a text source, predicting the most likely occurrence of words using its context

The observation probability of a word sequence $\mathbf{w} = \langle w_1 \dots w_m \rangle$:

$$P(\mathbf{w}) = P(w_1) \cdot \prod_{i=2}^m P(w_i | w_1 \dots w_{i-1}) \quad \text{with } w_i \in \Sigma$$

where $P(w_i | w_1 \dots w_{i-1})$ is the probability of having a word w_i , given a previous word history $w_1 \dots w_{i-1}$

The estimation of $P(\mathbf{w})$ is prohibitively expensive when vocabulary $|\Sigma|$ size becomes huge: $|\Sigma|^{i-1}$ different possible histories



- ▶ Assign a probability score to every word sequence
- ▶ Reflect a previous knowledge of a text source, predicting the most likely occurrence of words using its context

The observation probability of a word sequence $\mathbf{w} = \langle w_1 \dots w_m \rangle$:

$$P(\mathbf{w}) = P(w_1) \cdot \prod_{i=2}^m P(w_i | w_1 \dots w_{i-1}) \quad \text{with } w_i \in \Sigma$$

where $P(w_i | w_1 \dots w_{i-1})$ is the probability of having a word w_i , given a previous word history $w_1 \dots w_{i-1}$

The estimation of $P(\mathbf{w})$ is prohibitively expensive when vocabulary $|\Sigma|$ size becomes huge: $|\Sigma|^{i-1}$ different possible histories



- ▶ Assign a probability score to every word sequence
- ▶ Reflect a previous knowledge of a text source, predicting the most likely occurrence of words using its context

The observation probability of a word sequence $\mathbf{w} = \langle w_1 \dots w_m \rangle$:

$$P(\mathbf{w}) = P(w_1) \cdot \prod_{i=2}^m P(w_i | w_1 \dots w_{i-1}) \quad \text{with } w_i \in \Sigma$$

where $P(w_i | w_1 \dots w_{i-1})$ is the probability of having a word w_i , given a previous word history $w_1 \dots w_{i-1}$

The estimation of $P(\mathbf{w})$ is prohibitively expensive when vocabulary $|\Sigma|$ size becomes huge: $|\Sigma|^{i-1}$ different possible histories



$P(\mathbf{w})$ can be approximated by:

$$P(\mathbf{w}) \approx \prod_{i=1}^m P(w_i | \Phi_n(w_1 \dots w_{i-1})) = \prod_{i=1}^m P(w_i | w_{i-n+1} \dots w_{i-1})$$

The probability estimation of $P(w_i | w_{i-n+1} \dots w_{i-1})$ is usually computed from relative frequency counts $f(\cdot | \cdot)$:

$$P(w_i | w_{i-n+1} \dots w_{i-1}) = f(w_i | w_{i-n+1} \dots w_{i-1}) = \frac{C(w_{i-n+1} \dots w_{i-1} w_i)}{C(w_{i-n+1} \dots w_{i-1})}$$



$P(\mathbf{w})$ can be approximated by:

$$P(\mathbf{w}) \approx \prod_{i=1}^m P(w_i | \Phi_n(w_1 \dots w_{i-1})) = \prod_{i=1}^m P(w_i | w_{i-n+1} \dots w_{i-1})$$

The probability estimation of $P(w_i | w_{i-n+1} \dots w_{i-1})$ is usually computed from relative frequency counts $f(\cdot | \cdot)$:

$$P(w_i | w_{i-n+1} \dots w_{i-1}) = f(w_i | w_{i-n+1} \dots w_{i-1}) = \frac{C(w_{i-n+1} \dots w_{i-1} w_i)}{C(w_{i-n+1} \dots w_{i-1})}$$



Example

$S = \{(\text{el perro corre rapido}, 2),$
 $(\text{el tren azul corre veloz}, 2), (\text{el coche azul corre veloz}, 2)\}$

1-gram counts:

<s>	6
el	6
perro	2
corre	6
rapido	2
</s>	6
tren	2
azul	4
veloz	4
coche	2

2-gram counts:

<s> el	6
el perro	2
el tren	2
el coche	2
perro corre	2
corre rapido	2
corre veloz	4
rapido </s>	2
tren azul	2
azul corre	4
veloz </s>	4
coche azul	2

3-gram counts:

<s> el perro	2
<s> el tren	2
<s> el coche	2
el perro corre	2
el tren azul	2
el coche azul	2
perro corre rapido	2
corre rapido </s>	2
corre veloz </s>	4
tren azul corre	2
azul corre veloz	4
coche azul corre	2



- *n*-grams conditional probabilities can be estimated from raw text based on the relative frequency of word sequences

$$P(w_i | w_{i-n+1} \dots w_{i-1}) = \frac{C(qw_i)}{C(q)} = \frac{C(w_{i-n+1} \dots w_{i-1} w_i)}{C(w_{i-n+1} \dots w_{i-1})}$$

- An *n*-gram LM is said to be complete if all possible word sequences are represented ($|\Sigma|^n$) with an adequate estimation of their probabilities
- **Problem:** there are events (i.e. word sequences) that hardly ever or never occur in training → *Smoothing or discounting methods*



- n -grams conditional probabilities can be estimated from raw text based on the relative frequency of word sequences

$$P(w_i | w_{i-n+1} \dots w_{i-1}) = \frac{C(qw_i)}{C(q)} = \frac{C(w_{i-n+1} \dots w_{i-1} w_i)}{C(w_{i-n+1} \dots w_{i-1})}$$

- An n -gram LM is said to be complete if all possible word sequences are represented ($|\Sigma|^n$) with an adequate estimation of their probabilities
- **Problem:** there are events (i.e. word sequences) that hardly ever or never occur in training → *Smoothing or discounting methods*



- ▶ n -grams conditional probabilities can be estimated from raw text based on the relative frequency of word sequences

$$P(w_i | w_{i-n+1} \dots w_{i-1}) = \frac{C(qw_i)}{C(q)} = \frac{C(w_{i-n+1} \dots w_{i-1} w_i)}{C(w_{i-n+1} \dots w_{i-1})}$$

- ▶ An n -gram LM is said to be complete if all possible word sequences are represented ($|\Sigma|^n$) with an adequate estimation of their probabilities
- ▶ **Problem:** there are events (i.e. word sequences) that hardly ever or never occur in training → **Smoothing or discounting methods**



Smoothing Methods

- ▶ Back-Off
- ▶ Interpolation


Discounting Methods

- ▶ Good Turing
- ▶ Absolute Discounting
- ▶ Witten Bell
- ▶ Linear Discounting
- ▶ Kneser-Ney
- ▶ ...



$$P(w_i | w_{i-n+1}^{i-1}) = \begin{cases} P^*(w_i | w_{i-n+1}^{i-1}) & \text{if } C(w_{i-n+1}^i) > 0 \\ \beta(w_{i-n+1}^{i-1})P(w_i | w_{i-n+2}^{i-1}) & \text{otherwise} \end{cases}$$

where $P^*(\cdot|\cdot)$ is a discounted probability estimated to reserve mass for unseen events, and β is a back-off weight, which ensures the consistence of the model probabilities



Exercise 1. Compute the probability of the sentence “*el perro corre rapido*” with the 1-gram, the 2-gram, and the 3-gram that are introduced in the previous examples

Solution.

Probability with the 1-gram

$$\begin{aligned} p(\text{el perro corre rapido}) &= p(<s>) p(\text{el}) p(\text{perro}) p(\text{corre}) p(\text{rapido}) p(</s>) \\ &= \frac{6}{40} \frac{6}{40} \frac{2}{40} \frac{6}{40} \frac{2}{40} \frac{6}{40} = 1.27 \cdot 10^{-6} \end{aligned}$$

Probability with the 2-gram

$$\begin{aligned} p(\text{el perro corre rapido}) &= p(\text{el} | <s>) p(\text{perro} | \text{el}) p(\text{corre} | \text{perro}) p(\text{rapido} | \text{corre}) p(</s> | \text{rapido}) \\ &= \frac{6}{6} \frac{2}{6} \frac{2}{2} \frac{2}{6} \frac{2}{2} = 0.111 \end{aligned}$$

Probability with the 3-gram

$$\begin{aligned} p(\text{el perro corre rapido}) &= p(\text{perro} | <s>) p(\text{corre} | \text{el perro}) p(\text{rapido} | \text{perro corre}) \\ &= \frac{2}{6} \frac{2}{2} \frac{2}{2} \frac{2}{2} = 0.333 \end{aligned}$$



Applications of n-grams:

- ▶ Automatic speech recognition
- ▶ Machine translation
- ▶ Handwritten text recognition
- ▶ Spelling correction
- ▶ Text classification
- ▶ Plagiarism detection
- ▶ Language identification
- ▶ DNA sequence modeling
- ▶ ...



1. N-grams
2. Introduction to Probabilistic Finite-State Automata
3. Relation between N-grams and FSA
4. References



Introduction to Probabilistic Finite-State Automata

- ▶ **Free Monoid Σ^* :** Given a finite set Σ , Σ^+ is the set of all strings with finite length composed of elements from Σ

$$\Sigma^* = \Sigma^+ \cup \{\lambda\} \quad (\lambda \text{ is the empty string})$$

- ▶ **Grammar:** $G = (N, \Sigma, R, S)$

- N : Finite set of *non-terminal symbols*
- Σ : Finite set of *terminal symbols* or *primitives*
- $S \in N$: Initial non-terminal symbol or "*axiom*"
- $R \subset (N \cup \Sigma)^* N (N \cup \Sigma)^* \times (N \cup \Sigma)^*$: set of *rules* or *productions*

A rule is written as:

$$\alpha \rightarrow \beta, \quad \alpha \in (N \cup \Sigma)^* N (N \cup \Sigma)^*, \quad \beta \in (N \cup \Sigma)^*$$



► **Elemental derivation:** $\xRightarrow[G]{}$:

$$\mu \alpha \delta \xRightarrow[G]{} \mu \beta \delta \iff \exists (\alpha \rightarrow \beta) \in R, \mu, \delta \in (N \cup \Sigma)^*$$

► **Derivation** $\xRightarrow[G]{*}$:

It is a *finite sequence of elemental derivations*

A derivation d can be written as the corresponding sequence of rules of G

The *set of derivations* of $y \in \Sigma^*$ (such that $S \xRightarrow[G]{*} y$) is written as $D_G(y)$

A grammar G is *ambiguous* if $\exists y \in \Sigma^*$ such that $|D_G(y)| > 1$

► **Language generated by a grammar** G , $\mathcal{L}(G)$:

$$\mathcal{L}(G) = \left\{ y \in \Sigma^* \mid S \xRightarrow[G]{*} y \right\}$$



CHOMSKY HIERARCHY FOR RECURSIVE LANGUAGES

0: Non-restricted

1: Context-sensitive

$$\alpha \rightarrow \beta, \quad |\alpha| \leq |\beta|$$

2: Context-free

$$B \rightarrow \beta, \quad B \in N$$


3: **Regular or “Finite-state”**

$$A \rightarrow aB \text{ or } A \rightarrow a, \quad A, B \in N, \quad a \in \Sigma \cup \{\lambda\}$$



- ▶ **Regular grammars:** $G = (N, \Sigma, R, S)$,
Rules of R : $A \rightarrow aB \vee A \rightarrow a, A, B \in N, a \in \Sigma$
- ▶ **Finite-state automaton:** (non deterministic)
 $\mathcal{A} = (Q, \Sigma, \delta, q_0, F)$, $q_0 \in Q, F \subseteq Q, \delta : Q \times \Sigma \rightarrow 2^Q$
- ▶ **Equivalence:** For each regular grammar there exists a finite-state automaton that recognizes the same language^a

Example:

$$\begin{aligned}
 G &= (N, \Sigma, R, S); \\
 \Sigma &= \{a, b\}; N = \{S, A_1, A_2\}; \\
 R &= \{ S \rightarrow aA_1 \mid bA_2 \mid b, \\
 A_1 &\rightarrow aA_1 \mid bA_2 \mid b, \\
 A_2 &\rightarrow bA_2 \mid b \}
 \end{aligned}$$


$$\begin{aligned}
 \mathcal{A} &= \{Q, \Sigma, \delta, q_0, F\}; \\
 Q &= \{0, 1, 2\}, \\
 \Sigma &= \{a, b\}, \\
 q_0 &= 0, F = \{2\}
 \end{aligned}$$

$$\mathcal{L}(G) = \{b, ab, bb, aab, abb, bbb, \dots, aaabbbb, \dots\} = \mathcal{L}(\mathcal{A})$$

^a The reverse is not always true for stochastic languages!



Probabilistic Finite-state Automata (PFA): it is a tuple

$\mathcal{A} = \langle Q, \Sigma, \delta, I, F, P \rangle$, where:

- ▶ Q is a finite set of states
- ▶ Σ is the alphabet
- ▶ $\delta \subseteq Q \times \Sigma \times Q$ is a set of transitions
- ▶ $I : Q \rightarrow \mathbb{R}^{\geq 0}$ is the probability function of a state being an initial state
- ▶ $P : \delta \rightarrow \mathbb{R}^{\geq 0}$ is a probability function of transition between states
- ▶ $F : Q \rightarrow \mathbb{R}^{\geq 0}$ is the probability function of a state being a final state


I , P , and F are functions such that:

$$\sum_{i \in Q} I(i) = 1$$

$$\forall i \in Q, F(i) + \sum_{v \in \Sigma, j \in Q} P(i, v, j) = 1$$



Example:



Being $S_{\mathcal{A}}(y)$ all the state sequences in \mathcal{A} that generate y :

- **Probability of a string generated by a PFA \mathcal{A}**

$$p_{\mathcal{A}}(y) = \sum_{s \in S_{\mathcal{A}}(y)} p_{\mathcal{A}}(y, s) \rightarrow \text{forward algorithm}$$

- **Probability of the best sequence state for a string generated by a PFA \mathcal{A}**

$$p_{\mathcal{A}}(y) = \max_{s \in S_{\mathcal{A}}(y)} p_{\mathcal{A}}(y, s) \rightarrow \text{Viterbi algorithm}$$

Language defined by a PFA \mathcal{A}

$$\mathcal{L}(\mathcal{A}) = \{y \in \Sigma^* \mid p_{\mathcal{A}}(y) > 0\}$$





Relevant algorithms on PFA

- ▶ PFA operations are underpinned by an algebraic structure called **semiring**: two operations, two monoids (one of them commutative), one operation distributive with respect to the other operation
- ▶ PFA and Weighted Finite State Automaton are equally powerful
- ▶ PFA is a special case of Weighted Finite State Transducers (not output label)
- ▶ PFA operations: composition, determinization, epsilon-removal, weight-pushing, minimization, projection, shortest-path, concatenation, pruning, ...



PFA composition

Combine PFAs in such a way that the resulting PFA is the intersection of both languages




1. N-grams
2. Introduction to Probabilistic Finite-State Automata
3. Relation between N-grams and FSA
4. References




n -grams can be represented using deterministic PFA


Examples of PFA representing 1-grams, 2-grams and 3-grams:



1-grams



2-grams: $w_{i-1} w_i$




3-grams: $w_{i-2} w_{i-1} w_i$

For **2-grams**: $q = w_{i-1}$ and $q' = w_i$.


For **3-grams**: $q = w_{i-2} w_{i-1}$ and $q' = w_{i-1} w_i$.



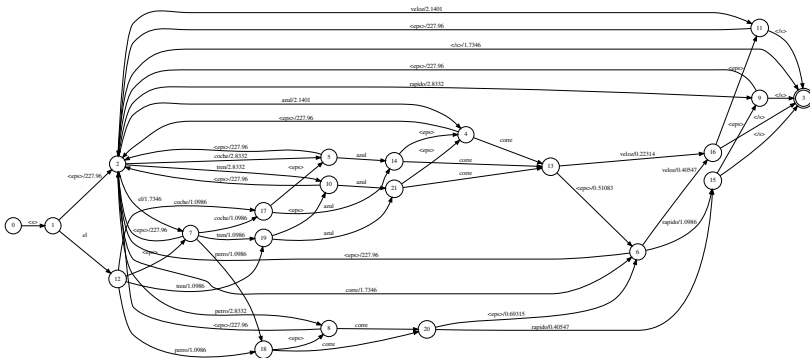
Example: 1-gram




Example: 2-gram




Example: 3-gram



Consequences of back-off as FSA: non-aggressive prune (2-gram)



Consequences of back-off as FSA: aggressive prune (2-gram)



1. N-grams
2. Introduction to Probabilistic Finite-State Automata
3. Relation between N-grams and FSA
4. References



- ▶ P. Dupont, F. Denis, and Y. Esposito. *Links between probabilistic automata and hidden Markov models: probability distributions, learning models and induction algorithms*, Pattern Recognition, 38:1349–1371, 2005.
- ▶ J.E. Hopcroft and J.D. Ullman. *Introduction to Automata Theory, Languages and Computation*. Addison-Wesley, 1979.
- ▶ F. Jelinek, *Statistical Methods for Speech Recognition*, MIT Press, 1998.
- ▶ R.A. Thompson *Determination of probabilistic grammars for functionally specified probability-measure languages*. IEEE Transactions of Computers, c-23(6):603–614, 1974.
- ▶ E. Vidal, F. Thollard, C. de la Higuera, F. Casacuberta, R. Carrasco, *Probabilistic finite-state machines - Part I*, IEEE Transactions on Pattern Analysis Machine Intelligence 27 (7):1013–1039, 2005.
- ▶ D. Jurafsky and J.H. Martin. *Speech and Language Processing*, Chapter 4, 2014.

